



Distributional formal semantics

Noortje J. Venhuizen^{a,*}, Petra Hendriks^b, Matthew W. Crocker^a,
Harm Brouwer^a

^a Saarland University, Department of Language Science and Technology, 66123 Saarbrücken, Germany

^b University of Groningen, Center for Language and Cognition Groningen (CLCG), P.O. Box 716, 9700AS Groningen, The Netherlands

ARTICLE INFO

Article history:

Available online 5 May 2021

Keywords:

Formal semantics
Distributional semantics
Compositionality
Probability
Inference
Incrementality

ABSTRACT

Natural language semantics has recently sought to combine the complementary strengths of formal and distributional approaches to meaning. However, given the fundamentally different ‘representational currency’ underlying these approaches—models of the world versus linguistic co-occurrence—their unification has proven extremely difficult. Here, we define Distributional Formal Semantics, which integrates distributionality into a formal semantic system on the level of formal models. This approach offers probabilistic, distributed meaning representations that are inherently compositional, and that naturally capture fundamental semantic notions such as quantification and entailment. Furthermore, we show how the probabilistic nature of these representations allows for probabilistic inference, and how the information-theoretic notion of “information” (measured in Entropy and Surprisal) naturally follows from it. Finally, we illustrate how meaning representations can be derived incrementally from linguistic input using a recurrent neural network model, and how the resultant incremental semantic construction procedure intuitively captures key semantic phenomena, including negation, presupposition, and anaphoricity.

© 2021 The Authors. Published by Elsevier Inc. This is an open access article under the CC BY-NC-ND license (<http://creativecommons.org/licenses/by-nc-nd/4.0/>).

1. Introduction

Traditional formal approaches to natural language semantics capture the meaning of linguistic expressions in terms of their logical interpretation within abstract formal models. Central to these approaches—which range from first-order predicate logic [1] to event-based [2] and dynamic semantic approaches, such as Discourse Representation Theory [3]—are the notions of entailment and compositionality, which describe how meanings are related to each other (entailment) and how they can be combined to form complex meanings (compositionality). An alternative approach to natural language semantics, which has recently gained much interest, is distributional semantics. This approach characterizes the meaning of lexical items in a usage-based manner; namely as distributional vectors that capture the co-occurrences between words [4–7]. The main advantage of such approaches is that the distributional representations inherently encode semantic similarity and relatedness between lexical items [8], and that they can be derived empirically from language data. It has, however, proven extremely difficult to incorporate the traditional semantic notions of entailment and compositionality within such a distributional semantic framework [9].

Indeed, while formal semantics focuses on proposition-level (sentence) meanings and distributional semantics focuses on the level of words, there has been considerable interest in bringing together the strengths of both approaches within a

* Corresponding author.

E-mail address: noortjev@coli.uni-saarland.de (N.J. Venhuizen).

single formalism. This has been attempted, for instance, by defining a notion of composition on top of the distributional representations, using vector operations [10], or by using more complex structures (e.g., matrices and tensors) in addition to vectors to represent lexical expressions [11–17]. Although this has been shown to produce interesting results when applied to adjective-noun modification [18], the approach has difficulties in representing the meaning of complex, multi-argument expressions (see [19,7] for reviews). Conversely, it has been attempted to incorporate distributional semantic representations within the logical forms that derive from formal models of meaning [20–22]. This approach aims to exploit the complementary strengths of formal and distributional semantics: formal mathematical machinery to account for sentence-level inference combined with a distributional account for lexical-level similarity and entailment [23]. While approaches such as these do indeed go a long way in extending formal semantics with a distributional component, they do so in a non-integrative manner; that is, while distributional semantic representations can guide the construction of logical form (e.g., [21]), logical inference remains part and parcel of the formal semantic system itself. This is due to the fact that formal and distributional semantics operate on a fundamentally different ‘representational currency’: formal semantics defines propositional meaning in terms of models of the world, whereas distributional semantics defines lexical meaning in terms of linguistic context.

An alternative approach from cognitive science, however, extends the situation-state space framework from Golden and Rumelhart [24] into a vector space model that defines meaning in a distributed manner relative to states of the world [25]. Here, we generalize this approach to reveal its full logical capacity: we introduce a notion of distributionality into a formal semantic system by defining meaning in terms of co-occurrence across formal models. The resulting framework, called Distributional Formal Semantics (DFS; as introduced in [26]), defines a meaning space for describing propositional-level meaning in terms of co-occurrence across a set of logical models: individual models are treated as observations, or cues, for determining the truth conditions of logical expressions—analogue to how individual linguistic contexts are cues for determining lexical meaning in distributional semantics. Based on a set of propositions \mathcal{P} , we define a set of logical models $\mathcal{M}_{\mathcal{P}}$ that together reflect the state of the world both truth-conditionally and probabilistically (i.e., reflecting the probabilistic structure of the world). The meaning of a proposition is defined as a vector within the meaning space constituted by $\mathcal{M}_{\mathcal{P}}$, which reflects its truth or falsehood relative to each of the models in $\mathcal{M}_{\mathcal{P}}$. The meaning vectors capture the (probabilistic) truth conditions of individual propositions indirectly by defining the meaning of individual propositions in relation to all other propositions; propositions that have related meanings will be true in many of the same models, and hence have similar meaning vectors. In other words, the meaning of a proposition is defined in terms of the propositions that it co-occurs with—or, to paraphrase the distributional hypothesis formulated by Firth [27]: “You shall know a *proposition* by the company it keeps”.

In what follows, we will introduce the DFS framework and show how its meaning space offers compositional distributed representations that are inherently probabilistic and inferential (see also [26]), and allow for capturing the information-theoretic notions of Entropy and Surprisal (cf. [28,29]). We show how propositional, as well as sub-propositional meanings are represented in the DFS framework as points in the meaning space, and how the incremental derivation of utterance meaning can be modeled as a trajectory through the meaning space. To this end, we employ a recurrent network model that is trained to map utterances onto their corresponding semantics on a word-by-word basis. We show how each word induces a contextualized transition from one point in meaning space to another, and critically extend upon previous work by modeling the effect of various core semantic phenomena—negation, presupposition, anaphoricity, and quantification—on incremental semantic construction. Finally, we discuss how Distributional Formal Semantics offers a powerful synergy between formal and distributional approaches to meaning—which we argue to be complementary to the lexical semantic knowledge captured by ‘traditional’ distributional approaches—that allows for modeling natural language semantics from formal, empirical, and cognitive perspectives.

2. A framework for Distributional Formal Semantics

In Distributional Formal Semantics, meaning is defined relative to a (finite) set of formal models $\mathcal{M}_{\mathcal{P}}$, in which each model is defined in terms of a set of propositions \mathcal{P} . The propositions in \mathcal{P} may be simple (zero-ary predicates) or complex (predicates with multiple arguments). The models in $\mathcal{M}_{\mathcal{P}}$ are first-order models that can be represented as the set of propositions $P \subseteq \mathcal{P}$ that they satisfy. The set of models $\mathcal{M}_{\mathcal{P}}$, which can theoretically be viewed as a set of possible worlds in the tradition of Carnap [30], constitutes a meaning space relative to which propositional meaning (as well as sub-propositional meaning; see below) can be expressed. An example of the meaning space is shown in Fig. 1; here, the rows represent models as sets of (truth values of) propositions and the columns represent propositional meaning vectors. The meaning vector $\vec{v}(p)$ thus defines the meaning of proposition $p \in \mathcal{P}$ in terms of the models that satisfy p ; i.e., $\vec{v}(p)$ is the vector that assigns 1 to all $M \in \mathcal{M}_{\mathcal{P}}$ such that p is satisfied in M , and 0 otherwise:

$$\llbracket p \rrbracket^{\mathcal{M}_{\mathcal{P}}} = \vec{v}(p) \quad \text{s.t. for all } i \in \{1, \dots, |\mathcal{M}_{\mathcal{P}}|\} : \vec{v}_i(p) = 1 \text{ iff } M_i \models p \quad (1)$$

Since propositional meaning is defined directly at the level of model interpretation, propositions that are true in many of the same models will obtain similar meaning vectors. That is, in the meaning space that derives from $\mathcal{M}_{\mathcal{P}}$, propositional meaning is defined in terms of co-occurrence between propositions—rather than in terms of linguistic co-occurrence, as in Distributional Semantics (e.g., [8]).

		meaning vectors for propositions in \mathcal{P}				
		p_1	p_2	p_3	\dots	p_n
models $\mathcal{M}_{\mathcal{P}}$	M_1	1	0	0	\dots	1
	M_2	0	1	1	\dots	1
	M_3	1	1	0	\dots	0
	\dots	\cdot	\cdot	\cdot	\dots	\cdot
	M_m	0	1	0	\dots	0

Fig. 1. Example of a meaning space based on a set of models $\mathcal{M}_{\mathcal{P}} = \{M_1 \dots M_m\}$, with meaning vectors for the set of propositions $\mathcal{P} = \{p_1 \dots p_n\}$.

In order for the meaning vectors to correctly capture propositional meaning in terms of co-occurrence with other propositions, the meaning space should be structured in such a way that it reflects the structure of the world, both truth-conditionally and probabilistically. That is, in order to fully encode entailment between two arbitrary propositions p and q ($p \models q$), the set of models $\mathcal{M}_{\mathcal{P}}$ should be truth-conditionally complete, such that: $p \models q$ (for any $p, q \in \mathcal{P}$) iff all models in $\mathcal{M}_{\mathcal{P}}$ that satisfy p also satisfy q , and $p \models \neg q$ iff no model in $\mathcal{M}_{\mathcal{P}}$ that satisfies p also satisfies q . The meaning space thus inherently encodes truth-conditional constraints on propositional co-occurrence in terms of entailment; we will refer to such constraints as *hard* world knowledge constraints (see Table 2 in Section 3.1 below for examples of such hard world knowledge constraints). Moreover, *probabilistic* world knowledge is captured in the meaning space in terms of the fraction of the models in $\mathcal{M}_{\mathcal{P}}$ that satisfy a certain proposition—or combination of propositions—from \mathcal{P} . Hence, in order for the meaning space to reflect the structure of the world, its constituent set of models $\mathcal{M}_{\mathcal{P}}$ should capture the hard and probabilistic world knowledge constraints dictated by the world. We here present a top-down algorithm that induces a meaning space based on a high-level specification of the world (but see the discussion in Section 5.2 for an alternative, data-driven approach to defining a meaning space for DFS). Accordingly, we describe the formal properties of the resulting meaning space in terms of the compositionality of the meaning vectors, their probabilistic and inferential properties, and how it captures information-theoretic aspects of meaning (using the notions of Entropy and Surprisal).

2.1. Sampling the meaning space

Given a set of propositions \mathcal{P} , the goal is to sample a set of models $\mathcal{M}_{\mathcal{P}}$ that reflects hard and probabilistic constraints of the world (theoretically, there are $2^{\mathcal{P}}$ possible models, but many of these are ruled out by hard constraints on propositional co-occurrence). That is, each individual model $M \in \mathcal{M}_{\mathcal{P}}$ should capture the truth-conditional constraints on co-occurrence between the propositions in \mathcal{P} , and the full set of models $\mathcal{M}_{\mathcal{P}}$ should reflect any probabilistic constraints on the (co-)occurrence of the propositions in \mathcal{P} . Following [26], we employ an inference-driven, non-deterministic sampling algorithm that stochastically generates a set of models from a pre-defined set of propositions \mathcal{P} and a set of (hard and probabilistic) constraints on these propositions.

The models $M \in \mathcal{M}_{\mathcal{P}}$ that constitute the meaning space are defined as basic (first-order) models that can be represented as tuples $\langle U_M, V_M \rangle$ consisting of a universe of entities U_M and an interpretation function V_M that assigns (sets of) entities to the individual constants and predicates represented in \mathcal{P} . The set of models $\mathcal{M}_{\mathcal{P}}$ is sampled by incrementally constructing individual models in a stochastic, inference-driven manner, based on the closed-world assumption (i.e., the assumption that individual models describe full states-of-affairs in terms of the truth and falsehood of the propositions in \mathcal{P}). The universe U_M of each $M \in \mathcal{M}_{\mathcal{P}}$ is defined as a set of entities $U_M = \{e_1 \dots e_n\}$, based on the set of individual constants ($c_1 \dots c_n$) that derive from \mathcal{P} . The interpretation function V_M , in turn, is initialized to map each constant onto a unique entity: $V_M(c_i) = e_i$. Next, we incrementally construct the model $\langle U_M, V_M \rangle$ based on the set of propositions \mathcal{P} , while taking into account the hard and probabilistic world knowledge constraints. Critically, since $\langle U_M, V_M \rangle$ only satisfies propositions that are assigned the truth value “true”, it does not distinguish between propositions that are assigned the truth value “false”, and those that are still undecided during incremental sampling. Therefore, in parallel to $\langle U_M, V_M \rangle$ —which we call the *Light World*¹ model (*LM*)—we construct a *Dark World* model $DM = \langle U_M, V'_M \rangle$ that satisfies all propositions that are assigned the truth value “false” (relative to *LM*). Indeed, this parallel model construction allows us to formulate not only the truth conditions of a particular expression using constraints on the Light World, but also its “falsehood conditions”, by evaluating the complement of the constraints relative to the Dark World. For instance, given a Light World constraint of the form $\forall x.R(x)$, any incrementally constructed Light World that does not satisfy all propositions of the form $R(x)$ will violate this constraint, even if the truth values of these propositions are still undecided. By introducing a Dark World, the falsehood of the constraint $\forall x.R(x)$ relative to the Light World can be proven by finding an entity in the Dark World for which $R(x)$ holds (hence formalizing the following logical equivalence: $\neg \forall x.R(x) \Leftrightarrow \exists x.\neg R(x)$). In other words, violation of the constraint can be explicitly verified by checking whether the complement of the constraint ($\exists x.\neg R(x)$) is satisfied relative to the Dark World. Table 1 shows the full set of rules that define the complement of a formula φ .

Using the Light World (*LM*) and Dark World (*DM*) models, the sampling algorithm incrementally constructs a model based on the set of propositions \mathcal{P} , a set of probabilistic constraints on \mathcal{P} —i.e., the function $Pr(p, M)$ that returns a prob-

¹ Cf. The Legend of Zelda: A Link to the Past (Nintendo, 1992).

Table 1The complement $\bar{\varphi}$ of any well-formed formula φ is defined based on the following rules^a.

$\overline{\neg\varphi}$	=	$\neg\bar{\varphi}$	$\overline{\varphi \oplus \psi}$	=	$(\bar{\varphi} \wedge \bar{\psi}) \vee (\neg\bar{\varphi} \wedge \neg\bar{\psi})$	$\overline{\exists x.\varphi}$	=	$\forall x.\bar{\varphi}$
$\overline{\varphi \wedge \psi}$	=	$\bar{\varphi} \vee \bar{\psi}$	$\overline{\varphi \rightarrow \psi}$	=	$\neg\bar{\varphi} \wedge \bar{\psi}$	$\overline{\forall x.\varphi}$	=	$\exists x.\bar{\varphi}$
$\overline{\varphi \vee \psi}$	=	$\bar{\varphi} \wedge \bar{\psi}$	$\overline{\varphi \leftrightarrow \psi}$	=	$(\neg\bar{\varphi} \wedge \bar{\psi}) \vee (\bar{\varphi} \wedge \neg\bar{\psi})$	\bar{p}	=	p

^a Note that the complement of the implication was reported incorrectly in [26].**Algorithm 1** Sampling algorithm for deriving a DFS meaning space based on a set of propositions \mathcal{P} , a set of probabilistic constraints $\{Pr(p, M) | p \in \mathcal{P}\}$, and a set of hard world knowledge constraints \mathcal{C} .

1. Let $\langle LM, DM \rangle$ be the tuple consisting of the initialized Light World model $LM = \langle U_M, LV_M \rangle$ and the initialized Dark World model $DM = \langle U_M, DV_M \rangle$;
2. Verify whether the state of affairs $\langle LM, DM \rangle$ is consistent: $\langle LM, DM \rangle \models \mathcal{C}$ iff for all constraints $c \in \mathcal{C}$: either c is satisfied ($LM \models c$) or c is not falsified ($DM \models \bar{c}$). If $\langle LM, DM \rangle \not\models \mathcal{C}$, start again from step 1.
3. Select a random proposition $p \in \mathcal{P}$, such that $LM \not\models p$ and $DM \not\models p$;
4. Let LM_p be an extension of LM such that $LM_p \models p$, and let DM_p be an extension of DM such that $DM_p \models p$;
5. Check the consistency of the state of affairs in which p is satisfied in the Light World ($\langle LM_p, DM \rangle$), and the state of affairs in which p is satisfied in the Dark World ($\langle LM, DM_p \rangle$).
 - If $\langle LM_p, DM \rangle \models \mathcal{C}$ and $\langle LM, DM_p \rangle \models \mathcal{C}$, the truth/falsehood of p is determined probabilistically: let LM_p be the new Light World LM with probability $Pr(p, LM)$; otherwise, let DM_p be the new Dark World DM .
 - If $\langle LM_p, DM \rangle \models \mathcal{C}$ and $\langle LM, DM_p \rangle \not\models \mathcal{C}$, p is inferred to be true: let LM_p be the new Light World LM .
 - If $\langle LM_p, DM \rangle \not\models \mathcal{C}$ and $\langle LM, DM_p \rangle \models \mathcal{C}$, p is inferred to be false: let DM_p be the new Dark World DM .
 - If $\langle LM_p, DM \rangle \not\models \mathcal{C}$ and $\langle LM, DM_p \rangle \not\models \mathcal{C}$, the state of affairs is inconsistent: start over from step 1.
6. Repeat from step 2 until all propositions in \mathcal{P} are satisfied in either LM or DM ;
7. If $\langle LM, DM \rangle \models \mathcal{C}$, the final LM is stored as a sampled model.

ability for proposition p based on model M (see Appendix A)—and a set of hard world knowledge constraints \mathcal{C} . This sampling procedure is described in Algorithm 1 below (cf. [26]). This sampling algorithm uses the Light World and the Dark World to determine for a randomly selected proposition p whether its truth (or falsehood) violates any world knowledge constraints. If p is consistent with respect to both the Light World and the Dark World, it is determined probabilistically whether the proposition is “true” or “false” (with respect to the Light World). If, on the other hand, p is only consistent with respect to the Light World, it is inferred to be “true” with respect to the Light World. Conversely, if p is only consistent with respect to the Dark World, it is inferred to be “false” with respect to the Light World (i.e., p is satisfied by the Dark World). Finally, if p is neither consistent with respect to the Light World nor with respect to the Dark World, the current model is discarded. This state of affairs may arise from an interaction between individual hard world knowledge constraints, e.g., when adding p to the Light World violates one constraint and adding p to the Dark World satisfies the complement of another. This sampling procedure is repeated until the truth or falsehood of each proposition $p \in \mathcal{P}$ is determined. The final Light World model $\langle U_M, LV_M \rangle$ will be added to the sampled set of models $\mathcal{M}_{\mathcal{P}}$. The full sampling algorithm described here is implemented as part of the software package `DFS-TOOLS`.²

2.2. Compositionality in DFS

The meaning vectors from the DFS meaning space are fully compositional at the propositional level. Since the meaning vector $\vec{v}(p)$ of a proposition $p \in \mathcal{P}$ is defined in terms of whether or not it is satisfied by each of the models $M \in \mathcal{M}_{\mathcal{P}}$, its negation can be defined straightforwardly as the complement of $\vec{v}(p)$, i.e., the vector that assigns 1 to all $M \in \mathcal{M}_{\mathcal{P}}$ such that p is not satisfied in M , and 0 otherwise:

$$\llbracket \neg p \rrbracket^{\mathcal{M}_{\mathcal{P}}} = \vec{v}(\neg p) \quad \text{s.t. for all } i \in \{1, \dots, |\mathcal{M}_{\mathcal{P}}|\} : \vec{v}_i(\neg p) = 1 \text{ iff } M_i \not\models p \quad (2)$$

Similarly, the meaning of the conjunction $p \wedge q$, for $p, q \in \mathcal{P}$, is defined as the meaning vector $\vec{v}(p \wedge q)$ that assigns 1 to all $M \in \mathcal{M}_{\mathcal{P}}$ that satisfy both p and q , and 0 otherwise:

$$\llbracket p \wedge q \rrbracket^{\mathcal{M}_{\mathcal{P}}} = \vec{v}(p \wedge q) \quad \text{s.t. for all } i \in \{1, \dots, |\mathcal{M}_{\mathcal{P}}|\} : \vec{v}_i(p \wedge q) = 1 \text{ iff } M_i \models p \text{ and } M_i \models q \quad (3)$$

Using the negation and conjunction operators, the meaning of any other logical combination of propositions in the semantic space can be defined, thus allowing for meaning vectors representing expressions of arbitrary logical complexity. Moreover, these operations also allow for the definition of quantification. Since \mathcal{P} fully describes the set of propositions expressed in $\mathcal{M}_{\mathcal{P}}$, the (combined) universe of $\mathcal{M}_{\mathcal{P}}$ ($U_{\mathcal{M}_{\mathcal{P}}} = \{u_1, \dots, u_n\}$) directly derives from \mathcal{P} . Universal and existential quantification, then, can be formalized by replacing the quantifier variable in the sub-formula with each of the entities in $U_{\mathcal{M}_{\mathcal{P}}}$, and combining them using conjunction and disjunction, respectively:

$$\llbracket \forall x\varphi \rrbracket^{\mathcal{M}_{\mathcal{P}}} = \vec{v}(\forall x\varphi) \quad \text{s.t. for all } i \in \{1, \dots, |\mathcal{M}_{\mathcal{P}}|\} : \vec{v}_i(\forall x\varphi) = 1 \text{ iff } M_i \models \varphi[x \setminus u_1] \wedge \dots \wedge \varphi[x \setminus u_n] \quad (4)$$

² DFS-TOOLS is available as open source under the Apache Licence, Version 2.0: <https://github.com/hbrouwer/dfs-tools>.

$$\llbracket \exists x \varphi \rrbracket^{\mathcal{M}_{\mathcal{P}}} = \vec{v}(\exists x \varphi) \quad \text{s.t. for all } i \in \{1, \dots, |\mathcal{M}_{\mathcal{P}}|\} : \vec{v}_i(\exists x \varphi) = 1 \text{ iff } M_i \models \varphi[x \setminus u_1] \vee \dots \vee \varphi[x \setminus u_n] \quad (5)$$

where $\varphi[x \setminus u]$ is defined as the formula φ with every instance of x replaced by u . This formalization of quantification in the meaning space constituted by $\mathcal{M}_{\mathcal{P}}$ is directly in line with traditional formalizations of quantification, in which an assignment function is used to substitute variables for elements from the model universe. The difference is that here variables are replaced by elements from the combined universe of all models in $\mathcal{M}_{\mathcal{P}}$. As a result, quantification is strictly defined with respect to the full set of models, rather than individual models: for instance, $\forall x R(x)$ is only true in those models that satisfy $R(x)$ for the full set of entities in $U_{\mathcal{M}_{\mathcal{P}}}$, not just for those entities that occur in the current model. This ensures that entailments are preserved across models, e.g., $\forall x.R(x)$ entails $R(e)$ for all entities e across all models in $\mathcal{M}_{\mathcal{P}}$.

2.3. Probability and inference in DFS

Propositional meaning vectors in the DFS meaning space are defined in terms of the models that satisfy a proposition. As a result, the meaning vectors are inherently probabilistic; that is, a proposition that is satisfied by a large number of models has a high probability, and vice versa. Formally, this means that the probability of an individual proposition is determined by the fraction of models in $\mathcal{M}_{\mathcal{P}}$ that satisfy this proposition (following [26]):

$$P(p) = \frac{|\{M \in \mathcal{M}_{\mathcal{P}} \mid M \models p\}|}{|\mathcal{M}_{\mathcal{P}}|} \quad \text{for } p \in \mathcal{P} \quad (6)$$

Given the compositional operations defined above, this means that the probability of any logical combination of propositions can be defined; for instance, the conjunctive probability of two propositions p and q is defined as the fraction of models that satisfy both propositions p and q :

$$P(p \wedge q) = \frac{|\{M \in \mathcal{M}_{\mathcal{P}} \mid M \models p \text{ and } M \models q\}|}{|\mathcal{M}_{\mathcal{P}}|} \quad \text{for } p, q \in \mathcal{P} \quad (7)$$

These definitions define probabilities over propositional-level meanings that are represented within the DFS meaning space as binary meaning vectors—reflecting truth and falsehood with respect to the models in $\mathcal{M}_{\mathcal{P}}$. Critically, however, the meaning space constituted by $\mathcal{M}_{\mathcal{P}}$ is continuous, which means that intermediate points in space—represented as real-valued vectors—constitute valid (sub-propositional) meanings with respect to $\mathcal{M}_{\mathcal{P}}$. Intuitively, each component i of a real-valued meaning vector $\vec{v}(a)$ can be interpreted using fuzzy logic as describing the ‘degree of truth’ of sub-propositional meaning a (e.g., representing a word or a sequence of words) relative to model $M_i \in \mathcal{M}_{\mathcal{P}}$. In other words, real-valued vectors constitute meanings that cannot be directly expressed as combinations of propositions, but rather reflect a degree of uncertainty regarding the propositional-level meanings. More formally, based on the set of models $\mathcal{M}_{\mathcal{P}}$, we can define the vector space $\mathbb{R}^{|\mathcal{M}_{\mathcal{P}}|}$ that contains all real-valued vectors within the dimensions of $\mathcal{M}_{\mathcal{P}}$. This vector space consists of both binary meaning vectors (constituting propositional meanings) and real-valued vectors (constituting sub-propositional meanings), that all carry their own probability. To capture the probabilistic properties of the vectors from the vector space $\mathbb{R}^{|\mathcal{M}_{\mathcal{P}}|}$, we extend the definitions for prior and conjunctive probabilities defined above (see Equations (6) and (7)) to account for real-valued vectors by calculating the average value of their components (following [31]). That is, given that the probability of propositions is defined in terms of the fraction of models that satisfy a proposition, the probability of a sub-propositional meaning—represented as a real-valued vector in which each component represents a fuzzy ‘degree of truth’—can be defined as the sum of its components divided by the number of models in $\mathcal{M}_{\mathcal{P}}$:

$$P(a) = \frac{1}{|\mathcal{M}_{\mathcal{P}}|} \sum_i \vec{v}_i(a) \quad \text{for } \vec{v}(a) \in \mathbb{R}^{|\mathcal{M}_{\mathcal{P}}|} \quad (8)$$

Given the fuzzy logic interpretation of the real-valued components of meaning vectors, the conjunction of two (distinct) points in meaning space a and b can standardly be defined using point-wise vector multiplication [31]. Importantly, since multiplying a real-valued vector with itself does not necessarily result in the original vector (in other words, the operation is a continuous t-norm that is non-idempotent), we define the conjunctive probability of a point with itself as its prior probability: $P(a \wedge a) = P(a)$. This ensures analogous behavior between the probabilities associated with propositional meanings (i.e., binary meaning vectors) and sub-propositional meanings (i.e., real-valued meaning vectors). The conjunctive probability of two arbitrary points in meaning space a and b (such that $a \neq b$) is then defined as follows:

$$P(a \wedge b) = \frac{1}{|\mathcal{M}_{\mathcal{P}}|} \sum_i \vec{v}_i(a) \vec{v}_i(b) \quad \text{for } \vec{v}(a), \vec{v}(b) \in \mathbb{R}^{|\mathcal{M}_{\mathcal{P}}|} \quad (9)$$

The definitions of prior and conjunctive probability given in Equations (8) and (9) extend the prior and conjunctive probabilities of propositions (Equations (6) and (7)) to real-valued meaning vectors in the vector space $\mathbb{R}^{|\mathcal{M}_{\mathcal{P}}|}$. More specifically, in the case of propositional meaning vectors, each model that satisfies a proposition p (or combination thereof)

contributes $\frac{1}{|\mathcal{M}|}$ of probability mass to $P(p)$. Similarly, in the case of real-valued meaning vectors, each model that satisfies sub-propositional meaning a in a fuzzy manner to degree d , contributes $\frac{d}{|\mathcal{M}|}$ of probability mass to $P(a)$.

Given the definitions for prior and conjunctive probability, we can calculate the conditional probability of any—propositional or sub-propositional—meaning a relative to any other—propositional or sub-propositional—meaning b in the meaning space:

$$P(a|b) = \frac{P(a \wedge b)}{P(b)} \quad (10)$$

Hence, the meaning of an arbitrary point in meaning space is inherently related—in terms of probabilistic co-occurrence—to any other point in meaning space. As a result, the meaning vectors inherently encode logical dependencies between propositions, and combinations thereof. We can therefore employ the conditional probability between meaning vectors to formally define entailment and probabilistic inference within the meaning space (see also [32]). That is, meaning vector a is entailed by meaning vector b ($b \models a$) if the conditional probability $P(a|b)$ equals 1 (i.e., if a and b reflect propositional meanings, this means that any model in $\mathcal{M}_{\mathcal{P}}$ that satisfies a also satisfies b). In order to quantify probabilistic inference of a given b , the prior probability of a needs to be taken into account: if the conditional probability $P(a|b)$ is higher than the prior probability $P(a)$, this means that a is positively inferred from b ; conversely, if $P(a|b) < P(a)$, a is negatively inferred from b . The following inference score (see [31]) quantifies this probabilistic inference on a range from +1 to −1.

$$\text{inference}(a, b) = \begin{cases} \frac{P(a|b) - P(a)}{1 - P(a)} & \text{if } P(a|b) > P(a) \\ \frac{P(a|b) - P(a)}{P(a)} & \text{otherwise} \end{cases} \quad (11)$$

An inference score of 1 indicates that meaning vector a is perfectly inferred from b (i.e., b entails a : $b \models a$), an inference score of −1 indicates that the negation of a is perfectly inferred from b ($b \models \neg a$), and any inference score in between these extremes reflects either positive ($\text{inference}(a, b) > 0$) or negative probabilistic inference ($\text{inference}(a, b) < 0$). It is important to note that the inference score itself does not define a probability: rather, it quantifies how much the truth of proposition b increases (or decreases) the probability of proposition a , relative to its prior probability. Hence, $\text{inference}(a, b) = 0$ means that the posterior probability $P(a|b)$ is equal to the prior probability $P(a)$; i.e., knowing b does not increase nor decrease the certainty in a .

As described above, (propositional) meaning in DFS is defined in terms of co-occurrence between propositions. Given that probabilities are defined for all points in meaning space, the inference score can be employed to describe the meaning of any real-valued vector $\vec{v}(a)$ relative to any other vector, and, in particular, relative to any of the propositions $p \in \mathcal{P}$ ($\text{inference}(p, a)$, which quantifies how much proposition p is inferred from the meaning vector constituted by a). Hence, all points in the meaning space—be it those constituted by a binary meaning vector representing propositional meaning, or by a real-valued vector representing sub-propositional meaning—are inherently probabilistic as well as inherently related to each other, as can be quantified using the inference score.

2.4. Quantifying information in DFS

The probabilistic nature of the meaning vectors in DFS also allows us to characterize points in meaning space using information theory, as proposed by Shannon [33]. That is, information theory defines the concept of “information” from a communicative perspective with respect to the set of possible messages that can be used to determine a particular state-of-affairs (e.g., there are six possible messages to describe the outcome of a roll of a die). This is mathematically captured by the notion of *Entropy*, which quantifies the amount of uncertainty in a given (communicative) state; i.e., states of high uncertainty (high Entropy) on average require more messages (information, in *bits*) to be resolved than states of low uncertainty. In terms of the meaning space, each point in space is defined relative to the set of models $\mathcal{M}_{\mathcal{P}}$ that constitutes the meaning space. These models are specified as maximally consistent sets of propositions, reflecting fully specified states of affairs. Moreover, they can themselves be represented as vectors in meaning space, namely as the vector $\vec{v}(M)$ that is defined as the conjunction of all propositions $p \in \mathcal{P}$ that are satisfied in M and the negations of all propositions $p' \in \mathcal{P}$ that are not satisfied in M . As a result, each point in meaning space inherently captures uncertainty about which fully specified state of affairs (i.e., which $M \in \mathcal{M}_{\mathcal{P}}$) is the case. That is, the closer a point is to a fully specified state of affairs, the less uncertainty it contains and hence the lower its information value. To quantify this notion of Entropy, [29] define a probability distribution over the set of meaning vectors that identify (unique) models in $\mathcal{M}_{\mathcal{P}}$, i.e., $\mathcal{V}_{\mathcal{M}_{\mathcal{P}}} = \{\vec{v}(M) \mid M \in \{M_i \mid M_i \in \mathcal{M}_{\mathcal{P}}\} \text{ and } \vec{v}_i(M) = 1 \text{ iff } M_i = M\}$. For a point in space a , Entropy is defined as follows:

$$H(a) = - \sum_{\vec{v}(M) \in \mathcal{V}_{\mathcal{M}_{\mathcal{P}}}} P(\vec{v}(M) | a) \log P(\vec{v}(M) | a) \quad (12)$$

Following this definition, Entropy is zero if the point in meaning space defined by a identifies a unique model. If, on the other hand, all models are equally likely (i.e., the probability distribution over all possible models is uniform), Entropy will be maximal.

Given this notion of Entropy, the logical and probabilistic properties of the DFS meaning space can be directly linked to the information-theoretic notion of uncertainty. Moreover, in the context of incremental language processing, a change in Entropy, typically a reduction, is taken to induce processing difficulty ([34]; see [29] for discussion). Similarly, processing effort has been linked to the notion of Surprisal, which quantifies the expectancy of words in context [35,36]; the less expected a word is in a given context, the higher its Surprisal, and hence the higher its processing effort. Following [28], Surprisal can be defined within the meaning space as reflecting the ‘expectedness’ of a transition between two points in meaning space, as triggered by an individual word or, more generally, a message. That is, for a given message $m_{a,b}$ that triggers a transition in meaning space from point a to b , Surprisal is high if point b is unexpected given point a , and low otherwise:

$$S(m_{a,b}) = -\log P(b|a) \quad (13)$$

In other words, the logarithm of the conditional probability between meaning vectors b and a is inversely proportional to the processing effort induced by the transition triggered by $m_{a,b}$. Indeed, this definition of Surprisal captures the “self-information” [33] of a transition in meaning space, and Entropy reflects the average Surprisal over all possible transitions from one point to the next within the meaning space.

The information-theoretic notions of Entropy and Surprisal thus directly derive from the probabilistic nature of the DFS meaning representations. As such, they constitute a direct link between formal theories of entailment and inference, and theories of incremental natural language processing. That is, incremental processing in the meaning space entails navigating the meaning space on a word-by-word basis, such that each word induces a transition from one point to the next. These intermediate meaning states represent the meaning that is constructed up to a given word, and derive from the mapping of sentences onto (propositional) meaning vectors in the meaning space. Although this mapping can in theory be formalized using a semantic interpretation function, e.g., using set-theoretic machinery (see section 5 for a preliminary discussion), it can also be approximated directly using a neural network model. The neural network approach has the advantage that intermediate meaning states not only capture probabilities deriving from the structure of the meaning space, but also those deriving from the (probabilistic) structure of the language (see [28]). Below, we show how such a model of incremental semantic construction captures semantic phenomena such as negation, quantification, presupposition and anaphoricity, and how these phenomena affect the incremental processing dynamics of the model.

3. A model of incremental semantic construction

Building a recurrent neural network model of incremental meaning construction involves three core steps. First, we need to construct an appropriate meaning space that allows for capturing the semantic phenomena of interest, in this case negation, presupposition, anaphoricity, and quantification. Secondly, we need to define a language \mathcal{L} that allows for describing (complex) situations in the meaning space, including expressions pertaining to the relevant phenomena. Finally, we need to train the neural network model to successfully map the utterances from \mathcal{L} onto their corresponding semantics on a word-by-word basis. Below, we will discuss each of these steps in detail.

3.1. Constructing a meaning space

The first step is to construct set of models $\mathcal{M}_{\mathcal{P}}$ that defines a meaning space, based on a confined set of propositions \mathcal{P} . For the current model, we construct a set of propositions using a set of predicates (*enter*(p,l), *call*(p,s), *arrive*(s), *order*(p,o), *bring*(s,o), *pay*(p)), which are combined with one or more of the following constants as arguments: persons ($p \in \{\text{mike, will, elli, nancy}\}$),³ places ($l \in \{\text{bar, restaurant}\}$), servers ($s \in \{\text{barman, waiter}\}$), and orders ($o_{\text{food}} \in \{\text{fries, salad}\}$; $o_{\text{drink}} \in \{\text{cola, water}\}$). In addition, each of the constants is associated with the one-place predicate “*referent*(r)”, which can be interpreted as introducing the referent into the current discourse context (cf. the universe of discourse referents, as formalized in Discourse Representation Theory [3]).

Based on the resulting set of propositions \mathcal{P} ($|\mathcal{P}| = 51$), a meaning space was constructed by sampling a set of 10K models $\mathcal{M}_{\mathcal{P}}$ (using the sampling algorithm described in Section 2.1, as implemented in DFS-TOOLS), while taking into account world knowledge in terms of hard and probabilistic constraints on propositional co-occurrence. Table 2 shows the hard constraints that are taken into account for sampling the meaning space; i.e., each model $M \in \mathcal{M}_{\mathcal{P}}$ satisfies all of these constraints. The first constraint ensures that each individual model can be interpreted unambiguously in terms of the state of affairs it describes; since no explicit temporal information is encoded in the predicates, multiple ‘*enter*’ events for the same person would obscure the event structure of individual models and therefore affect entailment and inferencing. Constraints 2–5 pose general constraints on event co-occurrence. Note that due to the probabilistic sampling strategy, individual models satisfy only a subset of the propositions in \mathcal{P} . Therefore, (non-)co-occurrences between individual propositions need to be defined explicitly and do not necessarily follow from combinations between other constraints (e.g., consider a model

³ In the language \mathcal{L} used in the current model (described in Section 3.2 below) “will” only occurs as a proper name, not as an auxiliary verb.

Table 2

Hard constraints implemented in construction of the meaning space. Based on the set of propositions \mathcal{P} (see text), the sampling algorithm samples a set of models $\mathcal{M}_{\mathcal{P}}$, such that each model $M \in \mathcal{M}_{\mathcal{P}}$ satisfies all of these constraints.

No.	Constraint	Description
1	$\forall x \forall y \forall z (\text{enter}(x, y) \wedge z \neq y \rightarrow \neg \text{enter}(x, z))$	A person can only enter a single place.
2	$\forall x \forall y_{f/d} \forall z_{f/d} (\text{order}(x, y) \wedge z \neq y \rightarrow \neg \text{order}(x, z))$	A person can only order a single type of food/drink.
3	$\forall x (\neg (\text{enter}(x, \text{bar}) \wedge \text{call}(x, \text{waiter})))$	Waiter cannot be called in bar.
4	$\forall x (\neg (\text{enter}(x, \text{restaurant}) \wedge \text{call}(x, \text{bartender})))$	Bartender cannot be called in restaurant.
5	$\forall x (\neg (\text{call}(x, \text{waiter}) \wedge \text{call}(x, \text{bartender})))$	A person can only call either waiter or bartender.
6	$\forall x \forall y ((\text{enter}(x, y) \wedge \text{pay}(x)) \rightarrow \exists z (\text{order}(x, z)))$	Entering and paying implies that something is ordered.
7	$\forall x \forall y ((\text{order}(x, y) \wedge \text{pay}(x)) \rightarrow \exists z (\text{bring}(z, y)))$	Ordering and paying implies that the order is brought.
8	$\forall P \forall x (P(x) \rightarrow \text{referent}(x))$	One-place predicates assert their argument.
9	$\forall R \forall x \forall y (R(x, y) \rightarrow (\text{referent}(x) \wedge \text{referent}(y)))$	Two-place predicates assert both arguments.

that only contains *call* and *order* predicates for a particular person: while it may satisfy constraints 1–4, it could violate constraint 5, and hence be invalid). In order to make sure individual models capture the canonical order of events, constraints 6–7 encode dependencies between event occurrences. Finally, constraints 8–9 make sure that all arguments of a predicate are instantiated as referents. This will allow us to model presupposition and anaphoricity.

In addition to the hard constraints, the set of models $\mathcal{M}_{\mathcal{P}}$ as a whole reflects the probabilistic structure identified by a set of probabilistic constraints; see Appendix A for the full set of probabilistic constraints used for sampling the current meaning space. As follows from the sampling algorithm defined in Section 2.1, the sampling probability of a proposition p is defined relative to a model M ($Pr(p, M)$). The intuition behind this is that the probability of p may depend on the propositions that are satisfied relative to the model M that describes the state-of-affairs constructed so far. For instance, the sampling probability of person p ordering food, $Pr(\text{order}(p, o_{\text{food}}), M)$, is low in case $\text{enter}(p, \text{bar})$ is satisfied in the current model M , while it is high in case $\text{enter}(p, \text{restaurant})$ is satisfied in M (if none of these constraints is satisfied in M , the proposition $\text{order}(p, o_{\text{food}})$ will be sampled using a base probability; see Appendix A). During sampling, the truth/falsehood of proposition p is determined based on the probability $Pr(p, M)$ only in case the truth/falsehood of p cannot be inferred (based on the Light World and Dark World constructed so far; see Section 2.1). Hence, as the sampling probabilities may interact with the hard world knowledge constraints described in Table 2, the observed probabilities in the sampled meaning space may only indirectly reflect these sampling probabilities (as will be illustrated below).

To render it feasible to approximate the resultant meaning representations using a neural network model, we reduced the set of 10K models to a representative set of 150 models, which captures the probabilistic and truth-conditional structure of the world (using the model selection algorithm described in Appendix B). To illustrate the richness of the inferences contained in the resulting 150-dimensional meaning space, Fig. 2 shows the by-proposition inferences represented in this space for a subset of the propositions.⁴ In this figure, bright green values indicate maximal inference (i.e., entailment) of proposition a given proposition b , and bright red values indicate maximal inference of $\neg a$ given proposition b . Every value in between reflects probabilistic inference. The green diagonal shows that each proposition entails itself. All other bright green and red values show (non-)co-occurrences between propositions that derive from the hard world knowledge constraints; for instance, $\text{order}(\text{mike}, \text{salad})$ entails $\text{referent}(\text{mike})$ and $\neg \text{order}(\text{mike}, \text{fries})$. Furthermore, the probabilistic inferences reflect the probabilistic world knowledge constraints described in Appendix A: for instance, the proposition $a = \text{order}(\text{mike}, \text{salad})$ is negatively inferred given $b = \text{enter}(\text{mike}, \text{bar})$ ($\text{inference}(a, b) = -0.39$) and positively inferred given $b' = \text{enter}(\text{mike}, \text{restaurant})$ ($\text{inference}(a, b') = 0.33$), hence following the pattern induced by the probabilistic constraints (see constraints 5 and 6 in Table A.4). The inference values shown in Fig. 2 thus quantify the structure of the meaning space in terms of the hard and probabilistic co-occurrence constraints between individual propositions, thereby providing a suitable means to compare different meaning spaces. In fact, the model selection algorithm described in Appendix B exploits exactly this information to derive a reduced set of models that approximates the structure of the original meaning space.

3.2. Mapping utterances onto semantics

Based on the meaning space derived from the set of models $\mathcal{M}_{\mathcal{P}}$, we define a language \mathcal{L} based on a set of words ($|W| = 30$)⁵ and a grammar that combines these words into utterances (consisting of one or two sentences) that describe situations (in terms of combinations of propositions) within the meaning space. Table 3 shows the utterances generated by the language \mathcal{L} and their associated semantics. Our grammar generates 278 unique utterances that describe 270 unique situations (due to the restricted nature of the meaning space, in which there exists only a single *bar/restaurant*, the definiteness of the determiner has no effect on the meaning of assertive sentences, thus resulting in overlapping semantics). The grammar generates three basic types of utterances: assertive simple sentences (e.g., “mike enters a restaurant”), negated simple sentences (e.g., “mike doesn’t enter a restaurant”) and assertive two-sentence utterances (e.g., “mike entered the bar

⁴ Please refer to the electronic version for color figures.

⁵ The negated auxiliary verb “didn’t” is considered a single word.




Fig. 2. By-proposition inferences in the meaning space for a subset of the propositions: all *referent* predicates together with all propositions pertaining to *mike* and those pertaining to the servers *waiter* and *bartender*. Bright green values indicate maximal inference of proposition *a* given proposition *b* (i.e., the entailment $b \models a$), white values indicate that there is no inference of proposition *a* given *b*, and bright red values indicate maximal inference of $\neg a$ given *b* (i.e., $b \models \neg a$); all intermediate values reflect positive (light green) or negative (light red) probabilistic inferences. (For interpretation of the colors in the figure(s), the reader is referred to the web version of this article.)

he ordered water”).⁶ In addition, all two-sentence utterances also occur as existentially quantified structures, which use the indefinite noun phrase “someone” (e.g., “someone entered the restaurant he ordered salad”). The semantics associated with assertive simple sentences are the individual propositions described by these sentences, and asserted two-sentence utterances describe conjunctions of propositions. The semantics of negated simple sentences, in turn, is defined as a conjunction between the negated proposition and the associated presupposed referents (i.e., those introduced using a name or a definite determiner)—since each proposition entails all its arguments as referents in the meaning space (see Section 3.1), this is not explicitly encoded in the semantics of assertive sentences. Finally, the utterances starting with “someone” obtain an existentially quantified version of their associated semantics. Given the interpretation of existential quantification in DFS (see Section 2.2), this results in a disjunctive semantics over all persons ($p \in \{\text{mike}, \text{will}, \text{elli}, \text{nancy}\}$).

Note that we here do not define a lexical semantics for individual words. We use the compositional machinery from DFS to combine propositional meanings into the meanings of entire utterances. To capture sub-propositional meaning, we exploit the continuous nature of the meaning space. That is, the meaning of a sub-propositional expression is a real-valued vector that defines a point in the meaning space, which is positioned in between those points that instantiate the propositional meanings that the expression pertains to (e.g., the meaning of “*mike*” will be expressed as the meaning vector that is positioned in between the propositional meanings that pertain to *mike*; *enter(mike,bar/restaurant)*, *call(mike,bartender/waiter)*, etc.). In contrast to traditional formal approaches, the DFS approach does not define an operation (such as function composition) that simply combines the sub-propositional meanings of two subsequent expressions. Rather, sequences of words $w_1 \dots w_n$ define a trajectory $\langle \vec{v}_1, \dots, \vec{v}_n \rangle$ through the meaning space, where each \vec{v}_i represents the (sub-propositional) meaning induced by the sequence of words $w_1 \dots w_i$; that is, each word w_i induces a meaning in the context of the meaning assigned to its preceding words $w_1 \dots w_{i-1}$. Sub-propositional meaning thus derives from the incremental, context-dependent mapping from word sequences onto (complex) propositional meanings. Below, we use a Simple Recurrent neural Network (SRN), as proposed by Elman [37], to approximate this mapping (see also [28,26]), and we show how the structure

⁶ Two-sentence utterances are concatenated without punctuation because in this small language it will have no effect on meaning construction.

Table 3

Utterances in language \mathcal{L} and their associated semantics. Each line contains two variants of the utterance (i.e., asserted/negated or individual/quantified) with differential semantics, indicated using a and b . Variables identify persons ($p \in \{\text{mike, will, elli, nancy}\}$), places ($l \in \{\text{bar, restaurant}\}$), servers ($s \in \{\text{barman, waiter}\}$), and orders ($o \in \{\text{cola, water, fries, salad}\}$). Utterances only describe possible situations in the meaning space and pronouns only occur with suitable antecedents ($\exists x_{m/f}$ restricts quantifier scope to male/female entities based on the pronoun).

Utterance _[a/b]	Semantics _a	Semantics _b
p [entered/didn't enter] a l	$\text{enter}(p,l)$	$\neg(\text{enter}(p,l)) \wedge \text{referent}(p)$
p [entered/didn't enter] the l	$\text{enter}(p,l)$	$\neg(\text{enter}(p,l)) \wedge \text{referent}(p) \wedge \text{referent}(l)$
p [called/didn't call] the s	$\text{call}(p,s)$	$\neg(\text{call}(p,s)) \wedge \text{referent}(p) \wedge \text{referent}(s)$
p [ordered/didn't order] o	$\text{order}(p,o)$	$\neg(\text{order}(p,o)) \wedge \text{referent}(p)$
p [paid/didn't pay]	$\text{pay}(p)$	$\neg(\text{pay}(p)) \wedge \text{referent}(p)$
the s [arrived/didn't arrive]	$\text{arrive}(s)$	$\neg(\text{arrive}(s)) \wedge \text{referent}(s)$
the s [brought/didn't bring] o	$\text{bring}(s,o)$	$\neg(\text{bring}(s,o)) \wedge \text{referent}(s)$
[p/someone] entered the l he/she ordered o	$\text{enter}(p,l) \wedge \text{order}(p,o)$	$\exists x_{m/f} (\text{enter}(x,l) \wedge \text{order}(x,o))$
[p/someone] entered the l he/she called the s *	$\text{enter}(p,l) \wedge \text{call}(p,s)$	$\exists x_{m/f} (\text{enter}(x,l) \wedge \text{call}(x,s))$
[p/someone] called the s he/she ordered o	$\text{call}(p,s) \wedge \text{order}(p,o)$	$\exists x_{m/f} (\text{call}(x,s) \wedge \text{order}(x,o))$
[p/someone] called the s he/she paid	$\text{call}(p,s) \wedge \text{pay}(p)$	$\exists x_{m/f} (\text{call}(x,s) \wedge \text{pay}(x))$
[p/someone] called the s he brought o	$\text{call}(p,s) \wedge \text{bring}(s,o)$	$\exists x (\text{call}(x,s) \wedge \text{bring}(s,o))$
[p/someone] called the s he arrived	$\text{call}(p,s) \wedge \text{arrive}(s)$	$\exists x (\text{call}(x,s) \wedge \text{arrive}(s))$

* Only allowing possible combinations of Location-Server: bar-bartender/restaurant-waiter.




Fig. 3. Simple Recurrent neural Network (SRN). Boxes represent groups of artificial neurons, and solid arrows between boxes represent full projections between the neurons in a projecting and a receiving group. The dashed lines indicate that the CONTEXT layer receives a copy of the activation pattern at the HIDDEN layer at the previous time-step. See text for details.

of the meaning space and the language \mathcal{L} combine in incremental processing. In particular, we examine how the processing of semantic phenomena such as negation, presupposition and anaphoricity emerges from this behavior.

3.3. Model specification

In order to derive a mapping from word sequences onto DFS meaning vectors, we train an SRN [37] to map sequences of words onto meaning vectors that represent propositional-level (i.e., utterance-final) meanings. The SRN consists of three groups of artificial logistic dot-product neurons: an input layer (30 units), a hidden layer (120), and an output layer (150) (see Fig. 3). Time in the model is discrete, and at each processing time-step t , activation flows from the input through the hidden layer to the output layer. In addition to the activation pattern at the input layer, the hidden layer also receives its own activation pattern at time-step $t - 1$ as input (effectuated through an additional context layer, which receives a copy of the activation pattern at the hidden layer prior to feedforward propagation). The hidden and the output layers both receive input from a bias unit (omitted in Fig. 3). We trained the model using bounded gradient descent [38] to map sequences of words onto a meaning vector representing the meaning of that utterance relative to the set of models $\mathcal{M}_{\mathcal{P}}$. More specifically, at each time step t during training, the model is presented with an individual localist word meaning representation at its input layer (a vector with a single hot bit),⁷ reflecting the current word in the unfolding utterance, and an utterance-final meaning vector at its output layer. At each time step t , the model thus effectively combines a word meaning representation (from the input layer) with an abstract representation of its context (reflected in the context layer) into a meaning vector in the DFS meaning space (at the output layer). Since individual words in context may map onto different utterance-final meanings with varying frequencies, at each processing step t the model will produce a vector at its

⁷ We use localist representations in order not to presuppose any word-internal structure. For more psychologically plausible representations, it is possible to employ representational schemes that encode phonetic, orthographic and/or semantic overlap between words.




Fig. 4. Visualization of the meaning space into three dimensions (using multidimensional scaling—MDS; see also Footnote 8) for a subset of the atomic propositions (those pertaining to *mike*, excluding *referent*), and the conjunctive meanings $\text{enter}(\text{mike}, \text{bar}) \wedge \text{order}(\text{mike}, \text{cola})$ and $\text{enter}(\text{mike}, \text{bar}) \wedge \text{order}(\text{mike}, \text{fries})$. Grey points represent (simple or complex) propositional meaning vectors. Colored points and lines show the word-by-word navigational trajectory of the model for the word sequence “mike entered [the] bar [he] ordered” (words in brackets are left out) with the continuations “cola” and “fries”. The numbers represent Surprisal values for the utterance-final words.

output layer that represents an abstraction over all possible utterance-final meanings. Critically, since the model is trained to map words onto meaning vectors, this output vector itself constitutes a point in meaning space (see Section 3.4 below).

Prior to training, the model’s weights were randomly initialized within the range of $(-.5, +.5)$. Each training item consisted of an utterance (a sequence of words represented by localist representations) and a meaning vector representing the utterance-final meaning. For each training item, error was backpropagated after each word, using a zero error radius of 0.05, meaning that no error was backpropagated if the error on a unit fell within this radius. Weight gradients were accumulated over epochs consisting of all training items. At the end of each epoch, weights were updated using a learning rate coefficient of 0.2 and a momentum coefficient of 0.9. Training lasted for 10000 epochs, after which the mean squared error was 0.79. The overall performance of the model was assessed by calculating the cosine similarity between each utterance-final output vector and each target vector for all utterances in the training data. After training, all output vectors had the highest cosine similarity to their own target (mean = .99; sd = .01), indicating that the model successfully learned to map utterances onto their corresponding semantics. We moreover computed how well the intended target could be inferred from the output of the model: $\text{inference}(\vec{v}_{\text{target}}, \vec{v}_{\text{output}})$. The average inference score over the entire training set was 0.88, which means that after processing an utterance, the model almost perfectly infers the intended meaning of the utterance.

3.4. Semantic construction as meaning space navigation

The model is trained to map each word in an utterance to the utterance-final meaning. As a single word sequence may occur in multiple items with different utterance-final meanings, the model learns to navigate the meaning space on a word-by-word basis. Fig. 4 provides a visualization of this navigation process. This figure is a three-dimensional representation of the 150-dimensional meaning space (for a subset of the atomic propositions and conjunctive meanings, derived using multidimensional scaling (MDS)).⁸ As this figure illustrates, meaning in DFS is defined in terms of co-occurrence; propositions that co-occur frequently in $\mathcal{M}_{\mathcal{P}}$ (e.g., $\text{enter}(\text{mike}, \text{restaurant})$ and $\text{call}(\text{mike}, \text{waiter})$) are positioned relatively close to each

⁸ Multidimensional scaling from 150 into 3 dimensions necessarily results in a significant loss of information. Therefore, distances between points in the meaning space shown in Fig. 4 should be interpreted with care. The grey points in this space correspond to propositional meaning vectors.




Fig. 5. Word-by-word inferences for the utterance “someone called the waiter she ordered cola” for a subset of the *referent* predicates. Inferences are calculated based on the word-by-word output of the model ($\vec{v}(w_t)$) and each of the propositions p , as follows: $\text{inference}(\vec{v}(p), \vec{v}(w_t))$. Identical line markers reflect arguments of the same type (person/location/server).

other (remember that the *waiter* can only be called in the *restaurant*). The colored points show the model’s word-by-word output for the word sequence “mike entered the bar he ordered”, with the continuations “cola” and “fries”. The navigational trajectory (indicated by the thick solid lines) illustrates how the model assigns intermediate points in meaning space to sub-propositional expressions, and approximates the utterance-final meaning at the final word.

Crucially, as shown in detail by [28], this trajectory is determined by the probabilistic structure of the meaning space (“world knowledge”) as well as the utterances on which the model was trained (“linguistic experience”). That is, at the word “mike”, the model navigates to a point in meaning space that is in between the meanings of the propositions pertaining to *mike*. Each consecutive word then results in a transition to a new point in space that best approximates the meaning up to that word. At “bar”, the model positions itself close to the meaning vector representing *enter(mike, bar)*, but does not fully commit to that meaning since its linguistic experience dictates that more input may come, which would result in a conjunctive utterance-final semantics. The effect of world knowledge, in turn, becomes clear at the final words. While the model was exposed to the utterances “mike entered the bar he ordered cola” and “mike entered the bar he ordered fries” equally often, the vector for “mike entered the bar he ordered” is closer to *enter(mike, bar) ∧ order(mike, cola)* than to *enter(mike, bar) ∧ order(mike, fries)*, because the former is more probable in the model’s knowledge of the world. This expectation is reflected in the Surprisal values (see Equation (13)) associated with the utterance-final words: “fries” is more surprising ($S(\text{fries}) = .96$) than “cola” ($S(\text{cola}) = .80$). For an elaborate investigation of the influence of world knowledge and linguistic experience on meaning space navigation and Surprisal, see [28].

4. Inference during semantic construction

The model maps utterances onto their corresponding semantics by navigating the meaning space on a word-by-word basis. Each word triggers a transition in meaning space from a point representing the meaning prior to encountering that word to a point that integrates it into the interpretation. Crucially, each transition reflects both the linguistic experience of the model, as well as by the world knowledge captured by the structure of the meaning space, and each point directly allows for probabilistic inferences about what is ‘understood’. That is, we can use the inference score described in Section 2.3 (see Equation (11)) to quantify how much a proposition p is inferred after processing word w_t : $\text{inference}(\vec{v}(p), \vec{v}(w_t))$.

Fig. 5 shows the word-by-word inferences for the utterance “someone called the waiter she ordered cola” for a selection of the *referent* propositions. This figure shows that at each word in the sentence, the model adjusts its inferences based on its linguistic experience and the structure of the meaning space. For instance, at the word “waiter”, the model perfectly infers that *referent(waiter)* is the case (blue line), and at the same time reduces its inference for *referent(bartender)* (orange line)—note that the hard world knowledge constraints do not exclude models that satisfy both *referent(waiter)* and *referent(bartender)*. Furthermore, at the sentence-final word “cola”, the model strongly infers *referent(restaurant)* (green line), while showing a negative inference for *referent(bar)* (red line), despite the fact that neither of these referents were mentioned in the sentence. This is an effect of the structure of the meaning space, in which *referent(waiter)* often co-occurs with *referent(restaurant)*, due to the hard world knowledge constraints (see Table 2). Finally, this figure shows that although the interpretation of the existentially quantified expression “someone” remains underspecified throughout the utterance, both the linguistic input (i.e., “she”) and the structure of the world (i.e., *elli* is more likely to order *cola* than *nancy*) guide the model toward an utterance-final interpretation in which *referent(elli)* (purple line) is inferred to a stronger degree than *referent(nancy)* (brown line).




Fig. 6. Processing of negation. The barplots show the inference scores for a relevant set of propositions given the model's output for the word sequences “will **ordered**” (left) and “will didnt **order**” (right)—critical words are shown in boldface.

Indeed, the ability to quantify the degree to which each given proposition is inferred at each point in the meaning space offers a powerful means to examine the dynamics of incremental semantic construction. In what follows, we will therefore harness probabilistic inference to investigate the construction dynamics of four key semantic phenomena: negation, presupposition, anaphoricity and quantification.

4.1. Negation

As described above, the meaning space inherently captures negation in terms of the models that do not satisfy (combinations of) propositions. This implies that each individual model in $\mathcal{M}_{\mathcal{P}}$ is interpreted as a full description of a state of affairs (rather than a partial observation, as in [25]), in which truth and falsehood directly reflect the state of the world. The negation semantics resulting from this meaning space were employed as the target semantics for the negated sentences presented to the model (see Table 3). Fig. 6 shows the contrast between an assertive (left) and a negated word sequence (right) for a selected set of inferences.

Both word sequences can be continued by introducing one of the possible orders (*cola*, *fries*, *salad*, *water*). At the word “ordered” in the utterance “will ordered” (left), the model has inferred that *referent(will)* is the case (blue bar). Moreover, the expected continuations are reflected by positive inferences for each of the propositions describing *will* ordering something (note that the personal preference for *water* over *cola* is also reflected in these expectations). By contrast, the negated word sequence “will didnt order” does not induce any of these probabilistic inferences; that is, although *referent(will)* is still inferred to be the case (see Section 4.2 below), there are no (positive or negative) inferences about any of the *order* propositions (small deviations from 0 are interpreted as noise resulting from the model's navigation through meaning space). This can be explained by the fact that the *order* predicates are partially mutually exclusive (*cola* excludes *water*, and *salad* excludes *fries*), and therefore all models $M \in \mathcal{M}_{\mathcal{P}}$ satisfy the negation of at least two of the *order* predicates. As a result, all models satisfy “will didnt order”, which means that all inferences approximate 0 (as the score $\text{inference}(a,b)$ quantifies the inference of *a* above and beyond its prior). Hence, negation—as part of the language presented to the model during training, as well as deriving from the structure of the meaning space—directly affects incremental semantic construction in terms of the online expectations represented in the model.

4.2. Presupposition

In order to capture (existential) presuppositions, the meaning space was constructed in such a way that all individual constants instantiated in a single model $M \in \mathcal{M}_{\mathcal{P}}$ were explicitly introduced using the *referent* predicate (similar to the instantiation of a universe of discourse referents in Discourse Representation Theory [3]; see also [39]). In the mapping of utterances onto semantics, these ‘existential’ predicates were explicitly incorporated to account for the imbalance between definite and indefinite noun phrases with regard to negation: whereas definite descriptions trigger existence presuppositions that ‘project’ from the scope of the negation (in line with [40]), indefinite descriptions generally do not trigger such presuppositions.⁹

Fig. 7 shows how these manipulations manifest in the online comprehension behavior of the model. The two panels at the top show that as part of an assertive sentence, definite and indefinite noun phrases result in the same inferences: the sentences “elli entered a restaurant/ elli entered the restaurant” both result in strong positive inferences (entailments) for *referent(elli)* (blue), *referent(restaurant)* (orange) and *enter(elli,restaurant)* (green). When embedded in a negated sentence, however, the inferences start to diverge. After processing the sentence “elli didnt enter a restaurant” (bottom left panel), the model infers that *referent(elli)* is the case and *enter(elli,restaurant)* is not the case. Critically, *referent(restaurant)* is also negatively inferred, indicating that the model finds itself in a state in meaning space in which *restaurant* is unlikely to be a referent (since $\neg \text{enter}(\text{elli}, \text{restaurant})$ can co-occur with other propositions that induce *referent(restaurant)*), this inference is

⁹ We here do not consider the use of ‘specific indefinites’, which do elicit such presuppositions in certain contexts (see, e.g., [41,42]).




Fig. 7. Processing of presupposition. The barplots show the inference scores for a relevant set of propositions given the model's output for the utterances "elli entered/didnt enter a **restaurant**" (top/bottom left) and "elli entered/didnt enter the **restaurant**" (top/bottom right)—critical words are shown in boldface.




Fig. 8. Processing of anaphora. The barplots show the inference scores for a relevant set of propositions given the model's output for the word sequence "mike called the bartender **he**" (left) and the continuations "mike called the bartender he **arrived**" (middle) and "mike called the bartender he **paid**" (right)—critical words are shown in boldface.

not maximally negative). The sentence "elli didnt enter the restaurant" (bottom right panel), on the other hand, does show a strong positive inference for *referent(restaurant)*, while at the same time inferring \neg *enter(elli, restaurant)*. This shows that the model is able to interpret presuppositions in context and to adjust its inferences accordingly.

4.3. Anaphoricity

The model was not explicitly trained to resolve the anaphoric dependencies of anaphoric expressions. Instead, anaphoric expressions (pronouns) were presented to the model as part of two-sentence utterances, which were associated with the appropriate (discourse-final) semantics. Crucially, pronouns were used to refer to both persons and servers (the latter always using the pronoun "he"), but their anaphoric antecedent was always disambiguated by either the prior context or the following sentence.

Fig. 8 shows an example of referential ambiguity of the pronoun. When processing the word "he" after the sentence "mike called the bartender" (left panel), the model shows moderately positive inferences for the propositions associated with possible continuations: both *arrive(bartender)* (orange) and *pay(mike)* (green) are positively inferred. Despite the fact that both continuations are equally likely based on the model's linguistic experience, *arrive(bartender)* is inferred to a slightly larger degree, which can be ascribed to a difference in conditional probabilities in the meaning space ($P(\text{arrive}(\text{bartender}) \mid \text{call}(\text{mike}, \text{bartender})) = .62$; $P(\text{pay}(\text{mike}) \mid \text{call}(\text{mike}, \text{bartender})) = .46$).¹⁰ When the model encounters the utterance-final word "arrived" (middle panel), the meaning is disambiguated to the vector representing $\text{call}(\text{mike}, \text{bartender}) \wedge \text{arrive}(\text{bartender})$ (because *arrive(mike)* is not a valid proposition in the meaning space). Instead, when the utterance is continued with "paid" (right panel), the model infers that *pay(mike)* is the case (again, *pay(bartender)* is not part of the meaning space), while still maintaining a slightly positive belief for *arrive(bartender)*.

Hence, after processing the ambiguous pronoun "he", the model shows positive inferences for both possible continuations, with a slight preference for one over the other. In other words, the inferences reflect the expectations of the model in terms of the possible continuations. We can use the notion of Surprisal (see Equation (13)) to quantify these expectations. Indeed, the Surprisal estimates reflect that the continuation "arrived" is more expected than "paid", because the former is less surprising ($S(\text{arrived}) = .44$) than the latter ($S(\text{paid}) = .66$), based on the probabilities of the meaning space and the language. In sum, the model entertains several possible interpretations upon encountering an ambiguous pronoun (some of

¹⁰ The model selection procedure in some cases results in (slight) probabilistic inferences that were not explicitly part of the probabilistic constraints of the world; in particular, reducing the number of models in $\mathcal{M}_{\mathcal{P}}$ may amplify co-occurrence probabilities between propositions.




Fig. 9. Processing of quantification. The barplots show the inference scores for a relevant set of propositions given the model's output for the word sequence "someone entered the **bar**" (left) and the continuations "someone entered the bar **he**" (middle) and "someone entered the bar **she**" (right)—critical words are shown in boldface.




Fig. 10. Processing of quantification (continued). The barplots show the inference scores for a relevant set of propositions given the model's output for the word sequence "someone entered the bar she **ordered**" (left) and the continuations "someone entered the bar she ordered **cola**" (middle) and "someone entered the bar she ordered **water**" (right)—critical words are shown in boldface.

which may be preferred over others), and incrementally uses the incoming linguistic information to resolve the pronoun and construct a coherent discourse-level meaning interpretation.

4.4. Quantification

Beyond assertive and negated sentences, the language \mathcal{L} also contains existentially quantified (two-sentence) utterances in which a quantified noun phrase ("someone") is used as the subject (see [26] for a similar model that also includes universally quantified utterances). The semantics for these utterances uses existential quantification in meaning space, which results in a disjunctive semantics over all possible subjects (i.e., the persons $p \in \{\text{elli}, \text{mike}, \text{nancy}, \text{will}\}$). From its linguistic experience, the model thus learns that "someone" can be used to refer to any of these subjects, without preferring any particular interpretation. Incremental navigation through meaning space, however, is driven by both the linguistic experience of the model as well as the structure of the meaning space. Given that the constituents of this disjunctive semantics may themselves differ in terms of their probability in meaning space, this will thus be reflected in the incremental inferences of quantified utterances. Moreover, since all quantified expressions presented to the model were part of a two-sentence utterance, the disjunctive semantics was in some cases limited to either male or female persons, depending on the pronoun used in the second sentence ("he/she").

The interaction between quantification and anaphoric processing is illustrated in Fig. 9. The left-most panel shows a selected set of relevant inferences for the sentence "someone entered the bar". At the word "bar", the model shows moderately positive inferences ($\sim .5$) for each of the predicates $\text{enter}(p, \text{bar})$. The inferences for $\text{enter}(\text{elli}, \text{bar})$ (blue) and $\text{enter}(\text{nancy}, \text{bar})$ (green) are slightly higher than the others, which is explained by the fact that they have a higher prior probability in the meaning space ($P(\text{enter}(\text{elli}, \text{bar})) = 0.29$; $P(\text{enter}(\text{nancy}, \text{bar})) = 0.25$; $P(\text{enter}(\text{mike}, \text{bar})) = 0.23$; $P(\text{enter}(\text{will}, \text{bar})) = 0.23$). The second and third panel show how these inferences change as the sentence is continued using either a male ("he") or a female pronoun ("she"). As expected, the gender of the pronoun determines which inferences are boosted; "he" results in higher inferences for $\text{enter}(\text{mike}, \text{bar})$ (orange) and $\text{enter}(\text{will}, \text{bar})$ (red), whereas "she" results in higher inferences for $\text{enter}(\text{elli}, \text{bar})$ and $\text{enter}(\text{nancy}, \text{bar})$. Moreover, the Surprisal estimates that derive from the model again reflect the model's expectations about possible continuations: the fact that the propositions describing females have a slightly increased inference score already at the word "bar" results in reduced Surprisal for "she" ($S(\text{she}) = 0.92$) relative to "he" ($S(\text{he}) = 1.12$).

To further illustrate how the model's inferences incrementally develop in the context of quantification, Fig. 10 shows how propositional inferences regarding *nancy* and *elli* develop after processing the utterance "someone entered the bar she ordered". At "ordered", the model is in a state of indecision, in which $\text{enter}(\text{elli}, \text{bar})$ (blue) and $\text{enter}(\text{nancy}, \text{bar})$ (orange) are both positively inferred. Moreover, for *elli* there is no preference between ordering *water* (green) and ordering *cola*

(red), whereas for *nancy* there seems to be a preference for *cola* (brown) as opposed to *water* (purple)—this is in line with the probabilistic constraints based on which the meaning space was created. When the utterance-final word “water” is processed (middle panel), the model maintains its disjunctive interpretation between $enter(elli, bar)$ and $enter(nancy, bar)$ and moreover boosts its inferences for $order(elli, water)$ and $order(nancy, water)$ to an equal degree (while reducing the inferences for ordering *cola*). Instead, when “cola” is encountered as utterance-final word (right panel), the model seems to show a slight preference for resolving its interpretation to *nancy*: $enter(nancy, bar)$ is boosted to a slightly higher degree than $enter(elli, bar)$, and $order(nancy, cola)$ obtains a higher inference than $order(elli, cola)$. This seems to suggest that after processing “cola”, the model is in a less uncertain state than after processing “water”, because there is a stronger preference for a particular interpretation of the pronoun. The uncertainty of a point in meaning space can be quantified using the notion of Entropy (see Equation (12)). Indeed, Entropy reveals that after “cola”, the model is in a slightly less uncertain state than after “water” ($H(\vec{v}(water)) = 3.60$; $H(\vec{v}(cola)) = 3.51$). Interestingly, the Surprisal estimates derived from the model seem to suggest that—in this particular example—confirming the expectation for $order(nancy, cola)$ is slightly less surprising ($S(cola) = 0.52$) than disconfirming it ($S(water) = 0.63$). For an investigation of the interaction between meaning-level Entropy and Surprisal during incremental meaning space navigation, see [29].

5. Discussion

Distributional Formal Semantics defines propositional and sub-propositional meaning within the meaning space, which is constituted from a set of propositions \mathcal{P} , and a set of formal models \mathcal{M} . On the propositional level, meaning vectors represent truth and falsehood with respect to the set of models \mathcal{M} , thereby inherently capturing co-occurrences between the propositions in \mathcal{P} . We have shown how sub-propositional meaning within the meaning space derives from the incremental navigation through meaning space, as modeled using a recurrent neural network. Crucially, we showed how semantic phenomena such as negation, presupposition, anaphoricity and quantification can be captured as part of the semantic construction procedure and how they affect meaning space navigation. Distributional Formal Semantics thus offers a powerful synergy between formal and distributional approaches that paves the way towards novel investigations into formal meaning representation and construction.

5.1. DFS and distributional semantics offer complementary meaning representations

Previous approaches that aimed to combine the strengths of formal and distributional semantics have tried to do so by either expanding distributional approaches to account for propositional-level inferences [13,15–17], or, conversely, by expanding formal approaches with a distributional component to account for lexical-level similarity [20–22]. By contrast, the DFS framework fundamentally integrates the distributional hypothesis into a formal semantic model, while maintaining the proposition-central perspective on meaning. That is, in terms of propositional-level meaning, the distributional hypothesis can be formulated as stating that propositions that occur in the same contexts—where contexts are models that represent states of the world, rather than linguistic contexts—represent similar meanings.

In DFS, the ‘representational currency’ is propositions, whereas in distributional semantics it is words. As a result, DFS allows us to model similarity at the propositional level; e.g., $call(mike, waiter)$ is more similar to $enter(mike, restaurant)$ than to $call(mike, bartender)$, since the former two propositions tend to co-occur in \mathcal{M} , whereas $call(mike, waiter)$ and $call(mike, bartender)$ never co-occur. Distributional semantics, on the other hand, models lexical similarity in terms of the distributional hypothesis [27]; “waiter” is similar to “bartender” as they occur in similar linguistic contexts. Crucially, the DFS approach and distributional semantics thus capture different notions of semantic similarity: while the latter offers representations that inherently encode feature-based lexical similarity between words, the former provides representations instantiating truth-conditional similarity between propositions. Hence, we argue for a division of labor between proposition-level and lexical distributional representations in which the former captures meaning in terms of the state of the world, and the latter captures the linguistic properties of lexical items (see also [43]). In other words, we take the DFS framework to be complementary to lexically-driven distributional semantics.

The complementary nature of these meaning representations is underlined by recent advances in the neurocognition of language, where evidence suggests that lexical retrieval (the mapping of words onto lexical semantics) and semantic integration (the integration of word meaning into the unfolding representation of propositional meaning) are two distinct processes involved in word-by-word sentence processing (see [44,45] for explicit neurocognitive models). More specifically, this neurocognitive perspective on comprehension suggests that there is no compositionality at the lexical level; that is, word forms in context are mapped into word meaning representations, which are subsequently integrated into a compositional phrasal-/utterance-level meaning representation. Indeed, this suggests that compositionality is only at play at the level of propositions, thus eschewing the need for compositionality at the lexical level.

5.2. Data-driven DFS

In DFS, the meaning of a proposition is captured by a meaning vector that defines for each model $M \in \mathcal{M}$ that constitutes the meaning space whether it satisfies the proposition or not. The quality of the meaning representations therefore critically depends on finding a set of models \mathcal{M} that truth-conditionally and probabilistically capture the structure of the world. In

Section 2 we described an algorithm to induce the meaning space from a high-level description of the structure of the world. An advantage of this approach is that it allows for defining the meaning space in a controlled manner, which is especially important when modeling experimental data. From a computational perspective, however, it would be interesting to investigate whether the meaning space—or the underlying world knowledge constraints—could be empirically derived from existing resources, such as knowledge bases and/or linguistic corpora.

There exists a wide variety of large-scale corpora that are annotated with sentence-level semantic information (for an overview, see [46]). For the purpose of generating a meaning space, the annotations should minimally reflect predicate-argument structure (e.g., based on PropBank [47], or FrameNet [48]). Crucially, however, a well-formed meaning space captures *all* co-occurrence probabilities between the propositions that constitute the space. This means that the resource should not only capture co-occurrence information between subsets of proposition-level meanings (e.g., those that are part of a single discourse, as in the Groningen Meaning Bank [49], or appear in a single textual entailment pair, as in the Stanford Natural Language Inference corpus [50]), but rather across the entire set of proposition-level meanings; that is, individual models in the DFS meaning space describe the truth-values of all propositions in \mathcal{P} , which would correspond to the union of the propositions occurring in all discourses, or across all textual entailment pairs in a corpus. Moreover, in order to maintain the distinction between probabilities deriving from the structure of the world, and those deriving from linguistic experience (see [28]), the co-occurrence probabilities represented in the meaning space should ideally derive from ‘world-knowledge’-driven inferences that are not confounded by linguistic co-occurrence (although in many cases these will of course align).

An example of a resource that captures such ‘world knowledge’-driven co-occurrences is the DeScript corpus [51], which contains crowd-sourced event sequence descriptions (ESDs) that describe typical, everyday activities (also called ‘scripts’ [52]). In addition, the DeScript corpus contains a gold standard alignment for a subset of the crowd-sourced ESDs. Although the ESDs are themselves linguistic, they contain ‘world-knowledge’-driven information regarding the events and that are typically associated with a particular script (e.g., baking a cake or fixing a bicycle), as well as the order in which they typically occur. Sampling a meaning space from such a corpus entails identifying a set of propositions \mathcal{P} (in this case, the events associated with a set of scripts), and collecting a large enough set of models \mathcal{M} that describe combinations between these propositions; whereas using individual ESDs may result in data sparseness, it is possible to sample models based on the co-occurrence information derived from the individual ESDs. In future work, we plan to evaluate this approach, and experiment with data-driven meaning space induction from different resources.

5.3. Set-theoretic meaning construction in DFS

The DFS meaning space is constituted by proposition-level meanings that are represented using binary meaning vectors. Since the meaning space itself is continuous, however, sub-propositional meaning can be captured using real-valued meaning vectors, which can be approximated using a Simple Recurrent neural Network. Within the neural network model, the meaning of a sub-propositional expression can be interpreted as the point in meaning space that is in between the (propositional) meaning vectors that represent the meanings of all possible continuations (see Fig. 4). In other words, the model navigates to a point in meaning space that is consistent with all interpretations that are captured by the current sub-propositional expression. An alternative approach to defining sub-propositional meaning is using sets of proposition-level meaning vectors (cf. [11]). That is, in line with the generalized quantifier approach that represents the meaning of a quantified expression as a set of propositions, the meaning of a sub-propositional expression in the meaning space can be represented as the set of propositional meaning vectors consistent with its meaning. For instance, the meaning of the proper name “mike” will be captured by the set of meaning vectors that describe the propositions that pertain to *mike*: $\llbracket \text{mike} \rrbracket_{\mathcal{S} \times \mathcal{P}} = \{\vec{v}(p) \mid p \in \{\text{enter}(\text{mike}, \text{bar}), \text{enter}(\text{mike}, \text{restaurant}), \text{call}(\text{mike}, \text{bartender}), \dots\}\}$. Since each meaning vector $\vec{v}(p)$ can be interpreted as describing the set of models that satisfy p , this set-theoretic interpretation effectively describes sub-propositional expressions as sets of sets of models.

Meaning construction, then, entails defining operations on sets of meaning vectors. In particular, the disjunction between two sets of meaning vectors can simply be defined as the union between these sets. Conjunction between sets of meaning vectors, in turn, means conjoining all elements of the first set of meaning vectors with all elements of the second set (based on the composition operation defined in Section 2.2). Since negating a set of meaning vectors implies that none of these is the case, the negation of a set of meaning vectors can be defined as the conjunctive closure over the negations of all individual meaning vectors. In order to capture the incrementality of meaning construction, a ‘merge’ operation (cf. merge between Discourse Representation Structures in DRT [3]) can be defined that asserts a set of propositional meaning vectors V into a context C —which is also defined as a set of meaning vectors—by selecting the subset of meaning vectors in V that is consistent with context C . A full description of set-theoretic interpretation in DFS is beyond the scope of the current manuscript, but an initial formalization of these operations can be found as part of DFS-TOOLS.

Crucially, sets of meaning vectors can be mapped back into the meaning space by finding the point in meaning space that is in between all elements of a given set. Formally, this means that the point in meaning space that corresponds to a set of meaning vectors is the (real-valued) vector that constitutes the arithmetic mean between these meaning vectors. In fact, given a completely balanced linguistic input, the recurrent neural network model is predicted to approximate exactly these intermediate points in space when capturing sub-propositional meaning. Indeed, it is the interaction between the structure of the meaning space and the structure of the linguistic experience that is difficult (if at all possible) to capture in

an ‘offline’ set-theoretic approach, while it is an inherent part of the way in which the neural network model maps language onto meaning.

5.4. DFS in cognitive models of language processing

DFS offers a powerful and flexible framework for modeling meaning and probabilistic inference in cognitive models of human language processing. For instance, a neural network model of language comprehension, similar to the one presented above, but employing meaning representations derived from an earlier formulation of the DFS framework (see [25]), has been used to successfully model the interaction between linguistic experience and world knowledge in comprehension [28]. Moreover, models employing such meaning representations have been shown to naturally capture inference and quantification [31], and generalize to unseen sentences and semantics, in both comprehension [31] and production [53]. Here, we have extended these results by showing how they capture phenomena such as negation, presupposition, and anaphoricity. Finally, as discussed above, the information-theoretic notions of Surprisal and Entropy directly derive from DFS representations (see [28,29]), thereby providing representation-grounded linking hypotheses between processing behavior in the model and behavioral correlates of human processing difficulty. Crucially, this spectrum of linking hypotheses can be extended even further by employing DFS representations in a neurocomputational model of the electrophysiology of language comprehension to also obtain direct estimates of neurophysiological processing, in particular, of the N400 and P600 components of the Event-Related brain Potential (ERP) signal; see [45].

Taken together, we believe that the power and flexibility of DFS representations, the incremental construction of these representations in neural network models, and the representation-grounded linking hypotheses to behavioral and neurophysiological indices of processing difficulty, provide a comprehensive workbench for 1) formal semantic theory grounded in psycholinguistic evidence, 2) offering foundations for psycholinguistic theory by formally explicating representations and mechanisms, and 3) overall integration of formal semantic and psycholinguistic approaches to the study of language. As such, we believe DFS has the potential to shed light on the processing nature, as well as the representational and mechanistic underpinnings thereof, of a vast spectrum of syntactic, semantic, and pragmatic phenomena.

6. Conclusion

We have proposed a framework for Distributional Formal Semantics (DFS), in which (sub-)propositional meaning is defined relative to a meaning space, which is constituted by a set of first-order models $\mathcal{M}_{\mathcal{P}}$ that is defined relative to a set of propositions \mathcal{P} (such that each model $M \in \mathcal{M}_{\mathcal{P}}$ can be defined in terms of the subset of propositions $P \subseteq \mathcal{P}$ that it satisfies). Within this meaning space, propositional meaning is represented by the meaning vector $\vec{v}(p)$ that represents the meaning of proposition $p \in \mathcal{P}$ in terms of the models that satisfy p ; i.e., $\vec{v}(p)$ is the vector that assigns 1 to all $M \in \mathcal{M}_{\mathcal{P}}$ such that p is satisfied in M . The meaning vector $\vec{v}(p)$ thus defines the meaning of p as a point in the meaning space. Within the meaning space, similar meanings (e.g., propositions with high co-occurrence) obtain similar meaning vectors that are positioned close to each other in space.

We have shown that the resultant distributed meaning representations are inherently compositional and probabilistic, and that they allow for capturing probabilistic inference and entailment. Moreover, we have shown how the information-theoretic notions of Surprisal (the ‘self-information’ of a transition in meaning space) and Entropy (the average Surprisal over all possible transitions from one point to the next) derive from these representations. Furthermore, sub-propositional meanings, which cannot be directly expressed as (combinations of) propositions, can be represented as real-valued vectors, constituting points in the meaning space that also capture their own probability. To derive these sub-propositional representations, we instantiated a semantic interpretation function that maps utterances onto DFS meaning representations. Rather than formalizing such a function using set-theoretic machinery, we have employed a recurrent neural network model to incrementally construct the semantics for an utterance by navigating the meaning space on a word-by-word basis. Crucially, we have shown that this model of incremental semantic construction naturally captures semantic phenomena such as negation, quantification, presupposition and anaphoricity, and how these phenomena affect the incremental processing dynamics of the model, as quantified by probabilistic inference, Surprisal, and Entropy.

DFS integrates the strengths of formal semantics and distributional semantics by incorporating a distributional component into a formal system. In DFS, like in formal semantics, the representational currency is propositions. In contrast to distributional semantics, which defines lexical meaning in terms of word co-occurrences, meaning in DFS is defined in terms of propositional co-occurrence, which reflects hard and probabilistic world knowledge constraints. As a result, we take DFS to be complimentary to distributional semantics; that is, where distributional semantics offers the representational machinery to capture lexical meaning, DFS offers utterance-level meaning representations. This distinction is in line with recent psycholinguistic theorizing, which suggests that compositionality may only be required at the utterance-level, thereby calling into question the need to directly incorporate aspects from formal semantics into distributional semantics.

In sum, DFS offers a powerful synthesis between formal and distributional approaches to semantics: distributed, utterance-level meaning representations that capture the similarity between (propositional) meanings, while strictly maintaining the formal properties of compositionality and entailment. As such, we believe that the DFS framework—implemented by DFS-TOOLS—paves the way towards novel investigations into the representation and construction of utterance-level meanings, and the relation between the formal, empirical and cognitive study of semantics.

Declaration of competing interest

None of the authors has any conflict of interest to disclose.

Acknowledgments

Funded by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) – Project-ID 232722074 – SFB 1102.

Appendix A. Probabilistic constraints

The inference-based sampling algorithm employs probabilistic constraints to determine the truth/falsehood of propositions that cannot be inferred, which is the case if a proposition is consistent both with respect to the Light World and with respect to the Dark World (see Algorithm 1 in Section 2.1). The set of probabilistic constraints used for sampling our meaning space is shown in Table A.4. During sampling, the truth of p is probabilistically determined relative to model M —the Light World—based on the probability described by $Pr(p, M)$, given that M satisfies pre-defined conditions (see the third column). The sampling algorithm employs these probabilistic constraints in an ordered manner such that the first matching constraint determines the probability. This means that if the model M that describes the state-of-affairs sampled so far does not satisfy any of the conditions defined for proposition p , p is sampled according to the base sampling probability (see constraint 19).¹¹ Note that a sampling probability of 1 (as in probabilistic constraints 14–17) means that whenever the truth of p is determined probabilistically, that is, if the truth/falsehood of p cannot be inferred, it will be assigned the truth value ‘true’ (with respect to the Light World). Conversely, a sampling probability of 0 (as in probabilistic constraint 18) means that p will never be probabilistically determined to be ‘true’—i.e., p will only be true relative to the Light World due to direct inference based on the hard world knowledge constraints (as described in Section 3.1). Hence, due to the inference-based nature of the sampling algorithm, the observed probabilities in the sampled meaning space may only indirectly reflect the sampling probabilities shown in Table A.4 (see Section 3.1 for discussion).

Table A.4

Probabilistic constraints used for sampling of propositions for the meaning space. The first column identifies the order and the second column provides a description of each of the probabilistic constraints. The sampling probability ($Pr(X, M)$), shown in column four, describes the sampling probability of proposition X relative to model M , given that M satisfies the conditions presented in column three. Individual variables correspond to logical constants of a particular type: persons ($p \in \{mike, will, elli, nancy\}$), places ($l \in \{bar, restaurant\}$), servers ($s \in \{barman, waiter\}$), and orders ($o_{food} \in \{fries, salad\}$; $o_{drink} \in \{cola, water\}$). The variables x and y are of the general type of entities, and X is of the type of propositions.

No.	Probabilistic constraint	Model conditions	Sampling probability
1	Persons tend to enter the same place	$M \models \exists x. enter(x, l)$	$Pr(enter(p, l), M) = 0.9$
2	Entering bar unlikely if person orders food	$M \models \exists o_{food}. order(p, o_{food})$	$Pr(enter(p, bar), M) = 0.1$
3	Entering bar likely if person orders drink	$M \models \exists o_{drink}. order(p, o_{drink})$	$Pr(enter(p, bar), M) = 0.9$
4	Entering rest. likely if person orders food	$M \models \exists o_{food}. order(p, o_{food})$	$Pr(enter(p, restaurant), M) = 0.9$
5	Low probability of ordering food in bar	$M \models enter(p, bar)$	$Pr(order(p, o_{food}), M) = 0.1$
6	High probability of ordering food in restaurant	$M \models enter(p, restaurant)$	$Pr(order(p, o_{food}), M) = 0.9$
7	Different persons unlikely to order the same order	$M \models \exists x. x \neq p \wedge order(x, o)$	$Pr(order(p, o), M) = 0.1$
8	Order unlikely to be brought if not ordered	$M \models \neg \exists x. order(x, o)$	$Pr(bring(s, o), M) = 0.1$
9	Low probability of paying if someone else pays in the same place	$M \models \exists x \exists y. enter(p, x) \wedge y \neq p \wedge enter(y, x) \wedge pay(y)$	$Pr(pay(p), M) = 0.1$
10	Personal preference (elli)	$M \models \top$	$Pr(order(elli, cola), M) = 0.9$
11	Personal preference (mike)	$M \models \top$	$Pr(order(mike, cola), M) = 0.9$
12	Personal preference (nancy)	$M \models \top$	$Pr(order(nancy, water), M) = 0.9$
13	Personal preference (will)	$M \models \top$	$Pr(order(will, water), M) = 0.9$
14	Bar presupposes bartender	$M \models referent(bar)$	$Pr(referent(bartender), M) = 1$
15	Restaurant presupposes waiter	$M \models referent(restaurant)$	$Pr(referent(waiter), M) = 1$
16	Bartender presupposes bar	$M \models referent(bartender)$	$Pr(referent(bar), M) = 1$
17	Waiter presupposes restaurant	$M \models referent(waiter)$	$Pr(referent(restaurant), M) = 1$
18	Base probability referents (infer only)	$M \models \top$	$Pr(referent(x), M) = 0$
19	Base probability propositions ¹¹	$M \models \top$	$Pr(X, M) = 0.6$

Appendix B. Model selection algorithm

We employ a model selection algorithm to reduce the dimensionality of the meaning space (following [28]). This means that based on a given set of models $\mathcal{M}_{\mathcal{P}}$, we define a new set of models $\mathcal{M}_{\mathcal{P}}^*$ such that: (i) $\mathcal{M}_{\mathcal{P}}^* \subset \mathcal{M}_{\mathcal{P}}$; (ii) $\mathcal{M}_{\mathcal{P}}^*$ captures

¹¹ The base sampling probability for propositions (see constraint 19) is set to 0.6 instead of a coin flip (0.5) in order to increase propositional co-occurrence across models.

Algorithm 2 Model selection algorithm that reduces a set of models $\mathcal{M}_{\mathcal{P}}$ to a subset $\mathcal{M}_{\mathcal{P}}^*$ of cardinality k , while preserving the structure of $\mathcal{M}_{\mathcal{P}}$ truth-conditionally and probabilistically.

1. Take a random proper subset of k models from $\mathcal{M}_{\mathcal{P}}$, and call this $\mathcal{M}_{\mathcal{P}}^*$;
2. Check if all propositional meaning vectors in $\mathcal{M}_{\mathcal{P}}^*$ are informative (i.e., there are no vectors that only contain zeros), otherwise return to step (1);
3. Compute an inference vector \vec{inf} containing the inference score $inf(a, b)$ for each combination of propositions $a, b \in \mathcal{P}$ for the original set of models ($\vec{inf}(\mathcal{M}_{\mathcal{P}})$) and the reduced set of models ($\vec{inf}(\mathcal{M}_{\mathcal{P}}^*)$);
4. Check if the reduced set of models encodes the same hard constraints as the original set of models (positive constraints: for all i , $\vec{inf}(\mathcal{M}_{\mathcal{P}}^*)(i) = 1$ iff $\vec{inf}(\mathcal{M}_{\mathcal{P}})(i) = 1$; negative constraints: for all i , $\vec{inf}(\mathcal{M}_{\mathcal{P}}^*)(i) = -1$ iff $\vec{inf}(\mathcal{M}_{\mathcal{P}})(i) = -1$), otherwise return to step (1);
5. Compute the similarity between $\mathcal{M}_{\mathcal{P}}^*$ and $\mathcal{M}_{\mathcal{P}}$ on the basis of the proposition-by-proposition inference scores in $\vec{inf}(\mathcal{M}_{\mathcal{P}}^*)$ and $\vec{inf}(\mathcal{M}_{\mathcal{P}})$ using Pearson's correlation coefficient $\rho(\vec{inf}(\mathcal{M}_{\mathcal{P}}^*), \vec{inf}(\mathcal{M}_{\mathcal{P}}))$;
6. If $\mathcal{M}_{\mathcal{P}}^*$ is the best approximation of $\mathcal{M}_{\mathcal{P}}$ so far ($\rho > \rho_{best}$), store it;
7. Start from step (1) and run the next iteration;
8. If we have reached X iterations, return the best $\mathcal{M}_{\mathcal{P}}^*$ found.

the same hard world knowledge constraints as $\mathcal{M}_{\mathcal{P}}$; and (iii) $\mathcal{M}_{\mathcal{P}}^*$ approximates the probabilistic inferences derived from $\mathcal{M}_{\mathcal{P}}$. Formally, we reduce the dimensionality of a meaning space by reducing the number of models in $\mathcal{M}_{\mathcal{P}}$ to k , such that $k < |\mathcal{M}_{\mathcal{P}}|$. The procedure, shown in Algorithm 2, is repeated for X iterations to arrive at a set of models $\mathcal{M}_{\mathcal{P}}^*$ (with $|\mathcal{M}_{\mathcal{P}}^*| = k$) that maximally reflects the knowledge encoded in the original space $\mathcal{M}_{\mathcal{P}}$. For the present model, we selected 150 models from the sampled set of 10K models (based on 50 iterations), resulting in a reduced meaning space with $\rho(\vec{inf}(\mathcal{M}_{\mathcal{P}}^*), \vec{inf}(\mathcal{M}_{\mathcal{P}})) = .91$.

References

- [1] G. Frege, *Begriffsschrift, eine der arithmetischen nachgebildete Formelsprache des reinen Denkens*, Nebert, 1879.
- [2] D. Davidson, Truth and meaning, *Synthese* 17 (1) (1976) 304–323.
- [3] H. Kamp, U. Reyle, *From Discourse to Logic: Introduction to Modeltheoretic Semantics of Natural Language, Formal Logic and Discourse Representation Theory*, Kluwer, Dordrecht, 1993.
- [4] P.D. Turney, P. Pantel, From frequency to meaning: vector space models of semantics, *J. Artif. Intell. Res.* 37 (2010) 141–188.
- [5] K. Erk, Vector space models of word meaning and phrase meaning: a survey, *Lang. Linguist. Compass* 6 (10) (2012) 635–653.
- [6] S. Clark, Vector space models of lexical meaning, in: S. Lappin, C. Fox (Eds.), *The Handbook of Contemporary Semantic Theory*, Wiley Online Library, 2015, pp. 493–522.
- [7] A. Lenci, Distributional models of word meaning, *Ann. Rev. Linguist.* 4 (2018) 151–171.
- [8] T.K. Landauer, S.T. Dumais, A solution to Plato's problem: the latent semantic analysis theory of acquisition, induction, and representation of knowledge, *Psychol. Rev.* 104 (2) (1997) 211–240.
- [9] G. Boleda, A. Herbelot, Formal distributional semantics: introduction to the special issue, *Comput. Linguist.* 42 (4) (2016) 619–635.
- [10] J. Mitchell, M. Lapata, Composition in distributional models of semantics, *Cogn. Sci.* 34 (8) (2010) 1388–1429.
- [11] M. Baroni, R. Zamparelli, Nouns are vectors, adjectives are matrices: representing adjective-noun constructions in semantic space, in: *Proceedings of the 2010 Conference on Empirical Methods in Natural Language Processing*, Association for Computational Linguistics, 2010, pp. 1183–1193.
- [12] F.M. Zanzotto, I. Korkontzelos, F. Fallucchi, S. Manandhar, Estimating linear models for compositional distributional semantics, in: *Proceedings of the 23rd International Conference on Computational Linguistics*, Association for Computational Linguistics, 2010, pp. 1263–1271.
- [13] M.S.B. Coecke, S. Clark, Mathematical foundations for a compositional distributed model of meaning, in: *Lambek Festschrift*, in: *Linguistic Analysis*, vol. 36, 2011, pp. 345–384.
- [14] R. Socher, B. Huval, C.D. Manning, A.Y. Ng, Semantic compositionality through recursive matrix-vector spaces, in: *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, Association for Computational Linguistics, 2012, pp. 1201–1211.
- [15] M. Baroni, R. Bernardi, R. Zamparelli, Frege in space: a program of compositional distributional semantics, *Linguist. Iss. Lang. Technol.* 9 (2014) 241–346.
- [16] E. Grefenstette, M. Sadrzadeh, Concrete models and empirical evaluations for the categorical compositional distributional model of meaning, *Comput. Linguist.* 41 (1) (2015) 71–118.
- [17] L. Rimell, J. Maillard, T. Polajnar, S. Clark, Relpron: a relative clause evaluation data set for compositional distributional semantics, *Comput. Linguist.* 42 (4) (2016) 661–701.
- [18] E.M. Vecchi, M. Marelli, R. Zamparelli, M. Baroni, Spicy adjectives and nominal donkeys: capturing semantic deviance using compositionality in distributional spaces, *Cogn. Sci.* 41 (1) (2017) 102–136.
- [19] W. Blacoe, M. Lapata, A comparison of vector-based representations for semantic composition, in: *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, 2012, pp. 546–556.
- [20] D. Garrette, K. Erk, R. Mooney, A formal approach to linking logical form and vector-space lexical semantics, in: *Computing Meaning*, Springer, 2014, pp. 27–48.
- [21] N. Asher, T. Van de Cruys, A. Bride, M. Abrusán, Integrating type theory and distributional semantics: a case study on adjective-noun compositions, *Comput. Linguist.* 42 (4) (2016) 703–725.
- [22] I. Beltagy, S. Roller, P. Cheng, K. Erk, R.J. Mooney, Representing meaning with a combination of logical and distributional models, *Comput. Linguist.* 42 (4) (2016) 763–808.
- [23] E. Chersoni, P. Blache, A. Lenci, Towards a distributional model of semantic complexity, in: *Proceedings of the Workshop on Computational Linguistics for Linguistic Complexity*, Association for Computational Linguistics, Stroudsburg, PA, 2016, pp. 168–177.
- [24] R.M. Golden, D.E. Rumelhart, A parallel distributed processing model of story comprehension and recall, *Discourse Process.* 16 (3) (1993) 203–237.
- [25] S.L. Frank, M. Koppin, L.G.M. Noordman, W. Vonk, Modeling knowledge-based inferences in story comprehension, *Cogn. Sci.* 27 (6) (2003) 875–910.
- [26] N.J. Venhuizen, P. Hendriks, M.W. Crocker, H. Brouwer, A framework for distributional formal semantics, in: R. Iemhoff, M. Moortgat, R. de Queiroz (Eds.), *Logic, Language, Information, and Computation*, Springer, Berlin, Heidelberg, 2019, pp. 633–646.
- [27] J.R. Firth, A synopsis of linguistic theory, 1930–1955, in: *Studies in Linguistic Analysis*, Philological Society, Oxford, 1957.
- [28] N.J. Venhuizen, M.W. Crocker, H. Brouwer, Expectation-based comprehension: modeling the interaction of world knowledge and linguistic experience, *Discourse Process.* 56 (3) (2019) 229–255, <https://doi.org/10.1080/0163853X.2018.1448677>.
- [29] N.J. Venhuizen, M.W. Crocker, H. Brouwer, Semantic entropy in language comprehension, *Entropy* 21 (12) (2019) 1159, <https://doi.org/10.3390/e21121159>.

- [30] R. Carnap, *Meaning and Necessity*, The University of Chicago Press, Chicago, 1947.
- [31] S.L. Frank, W.F.G. Haselager, I. van Rooij, Connectionist semantic systematicity, *Cognition* 110 (3) (2009) 358–379.
- [32] M. Richardson, P. Domingos, Markov logic networks, *Mach. Learn.* 62 (1–2) (2006) 107–136.
- [33] C.E. Shannon, A mathematical theory of communication, *Bell Syst. Tech. J.* 27 (3) (1948) 379–423.
- [34] J.T. Hale, Uncertainty about the rest of the sentence, *Cogn. Sci.* 30 (4) (2006) 643–672.
- [35] J.T. Hale, A probabilistic Earley parser as a psycholinguistic model, in: *Proceedings of the Second Meeting of the North American Chapter of the Association for Computational Linguistics on Language technologies*, Association for Computational Linguistics, Stroudsburg, PA, 2001, pp. 1–8.
- [36] R. Levy, Expectation-based syntactic comprehension, *Cognition* 106 (3) (2008) 1126–1177.
- [37] J.L. Elman, Finding structure in time, *Cogn. Sci.* 14 (2) (1990) 179–211.
- [38] D.L.T. Rohde, A connectionist model of sentence comprehension and production, Ph.D. thesis, Carnegie Mellon University, 2002.
- [39] N.J. Venhuizen, J. Bos, P. Hendriks, H. Brouwer, Discourse semantics with information structure, *J. Semant.* 35 (1) (2018) 127–169, <https://doi.org/10.1093/jos/ffx017>.
- [40] P.F. Strawson, On referring, *Mind* 59 (235) (1950) 320–344.
- [41] A. Kasher, D.M. Gabbay, On the semantics and pragmatics of specific and non-specific indefinite expressions, *Theor. Linguist.* 3 (1–3) (1976) 145–190.
- [42] B. Geurts, Specific indefinites, presupposition and scope, in: R. Bäuerle, U. Reyle, T.E. Zimmermann (Eds.), *Presuppositions and Discourse: Essays offered to Hans Kamp*, Emerald, Bingley, 2010, pp. 125–158.
- [43] K. Erk, What do you know about an alligator when you know the company it keeps?, *Semant. Pragmat.* 9 (17) (2016) 1–63, <https://doi.org/10.3765/sp.9.17>.
- [44] H. Brouwer, M.W. Crocker, N.J. Venhuizen, J.C.J. Hoeks, A neurocomputational model of the N400 and the P600 in language processing, *Cogn. Sci.* 41 (2017) 1318–1352, <https://doi.org/10.1111/cogs.12461>.
- [45] H. Brouwer, F. Delogu, N.J. Venhuizen, M.W. Crocker, Neurobehavioral correlates of surprisal in language comprehension: a neurocomputational model, *Front. Psychol.* (2021), <https://doi.org/10.3389/fpsyg.2021.615538>.
- [46] E.M. Bender, A. Lascarides, *Linguistic Fundamentals for Natural Language Processing II: 100 Essentials from Semantics and Pragmatics*, Synthesis Lectures on Human Language Technologies, vol. 12, Morgan & Claypool Publishers, 2019.
- [47] M. Palmer, D. Gildea, P. Kingsbury, The proposition bank: an annotated corpus of semantic roles, *Comput. Linguist.* 31 (1) (2005) 71–106.
- [48] C.F. Baker, C.J. Fillmore, J.B. Lowe, The Berkeley FrameNet project, in: *Proceedings of the 36th Annual Meeting of the Association for Computational Linguistics and 17th International Conference on Computational Linguistics*, Vol. 1, Association for Computational Linguistics, 1998, pp. 86–90.
- [49] J. Bos, V. Basile, K. Evang, N.J. Venhuizen, J. Bjerva, The Groningen meaning bank, in: N. Ide, J. Pustejovsky (Eds.), *Handbook of Linguistic Annotation*, Springer, Netherlands, Dordrecht, 2017, pp. 463–496.
- [50] S.R. Bowman, G. Angeli, C. Potts, C.D. Manning, A large annotated corpus for learning natural language inference, in: *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing, EMNLP*, Association for Computational Linguistics, 2015.
- [51] L.D. Wanzare, A. Zarcone, S. Thater, M. Pinkal, DeScript: a crowdsourced corpus for the acquisition of high-quality script knowledge, in: *The International Conference on Language Resources and Evaluation*, 2016.
- [52] R.C. Schank, R.P. Abelson, *Scripts, Plans, Goals, and Understanding: An Inquiry into Human Knowledge Structures*, Lawrence Erlbaum Associates, Hillsdale, NJ, 1977.
- [53] J. Calvillo, H. Brouwer, M.W. Crocker, Connectionist semantic systematicity in language production, in: A. Papafragou, D. Grodner, D. Mirman, J.C. Trueswell (Eds.), *Proceedings of the 38th Annual Conference of the Cognitive, Science Society*, Austin, TX, 2016, pp. 2555–3560.